

Deliverable D — The 30-Minute Crash Lesson (start here)

Welcome. By the end of this you will understand the single idea the whole unit is built on. Read it slowly, out loud if you can. Everything else you study will just be variations on what's here. No notes will be needed in the exam if this lives in your head.

1. What is ethics, really?

Ethics is the study of what we ought to do and what makes actions right or wrong. That's it. Not "what people happen to believe" but "what people *should* do."

There's a crucial split here:

- **Descriptive ethics** = describing what people actually believe ("most Australians think lying is wrong"). This is just reporting facts, like a survey.
- **Normative ethics** = working out what is *actually* right or wrong ("lying *is* wrong, and here's why"). This is the arguing part.

This unit is **normative**. We are not surveying opinions; we are arguing about what's true.

The single most important habit: an answer in ethics is never just a feeling. "I feel drone strikes are bad" scores nothing. You must turn the feeling into a **mini-argument**: a *claim* plus a *reason*. "Drone strikes are wrong *because* they kill civilians who pose no threat." The word "because" is your best friend. In your oral exam, every answer must contain a "because."

Think of it like maths homework: the teacher doesn't want just the answer, they want the working. In ethics, the reason *is* the working.

2. Why we set relativism aside

Relativism is the idea that there's no universal right or wrong — morality is just whatever each culture or person decides, so no one can ever be judged.

It sounds tolerant and open-minded. But we set it aside for a simple reason: if relativism were true, this entire unit would be pointless. If "the Nazis thought genocide was fine, so for them it was fine" is the end of the discussion, there's nothing left to study. We could never say a war crime was genuinely wrong — only that "we don't happen to like it."

We don't have to *prove* relativism false to move on. We just note that the unit assumes some moral claims are genuinely better-supported than others, and then we get to work arguing about which ones. Park relativism. Get on with the reasoning.

3. THE MASTER KEY: consequentialism vs deontology

This is the heart of everything. Almost every debate in the unit is a fight between these two families. Learn this and you can reconstruct half the course on the spot.

Consequentialism = only **outcomes** matter. An action is right if it leads to the best overall results — usually the most wellbeing, the fewest deaths, the most good on balance. The famous slogan: *the ends can justify the means*. A consequentialist adds up the consequences like a ledger and picks whatever number comes out best. (Utilitarianism is the most common version — "the greatest good for the greatest number.")

Everyday analogy: You break a promise to a friend to drive a stranger to hospital. A consequentialist says: obviously right — a life saved hugely outweighs a broken coffee date. Tally the results, pick the best total.

Deontology = some acts are right or wrong **in themselves**, whatever the outcome. There are moral *rules* and *duties* you must not break, even if breaking them would produce a better total. The slogan: *some things you just may not do*. (The strictest version is called **absolutism** — *never* break the rule, no matter what; for this unit you can treat "deontology" and "absolutism" as the same family, just remember absolutism is the no-exceptions end of it.)

Everyday analogy: A surgeon has five patients dying for want of organs and one healthy person in the waiting room. Killing the one to harvest organs for five would, on a pure tally, "save four net lives." Almost everyone recoils — *you may not butcher an innocent person, full stop, even for a great outcome*. That recoil is deontology talking. The person has a **right** not to be killed that "trumps" the good numbers.

Now the war example where they clash — burn this in:

It's wartime. Deliberately bombing an enemy city's civilians would, let's suppose, break enemy morale and end the war six months early, saving more lives overall than it costs.

- The **consequentialist** says: if the maths genuinely works — more lives saved than lost — then drop the bomb. Outcomes are all that matter.
- The **deontologist** says: no. Deliberately killing innocent civilians is murder, and murder is wrong in itself. You may not *deliberately target the innocent* even to win the war, even if the tally looks good. Their right not to be slaughtered trumps the calculation.

That single disagreement — *may you ever deliberately kill the innocent for a good result?* — is the engine of the entire course. Hiroshima, terrorism, drone strikes, killer robots: all of them are this same clash in new clothes.

4. The four supporting tools

These are refinements that mostly help the deontologist make their case precise. Each gets a tiny example.

(a) Intention vs foresight. There's a moral difference between *aiming at* a harm and merely *foreseeing* it as a side-effect.

You give a dying patient morphine *intending* to ease pain, foreseeing it may shorten their life — versus giving the same dose *intending* to kill them. Same death, different act. War version: bombing a weapons factory and foreseeing some neighbours will die is treated very differently from bombing the neighbours *in order to* terrorise. Deontologists think this distinction matters enormously.

(b) Doing vs allowing. There's a moral difference between *causing* a harm and *letting* it happen.

Pushing someone into a river versus walking past as they drown. Both end in a death; most people think actively drowning them is worse. War version: a country that *bombs* civilians seems to do something worse than one that merely *fails to prevent* a famine — even if the death tolls match.

(c) Rights and duties. A **right** is a protection a person has that can "trump" good outcomes; a **duty** is an obligation you owe them.

You have a right not to be tortured even if torturing you would entertain a stadium of thousands. The pleasure of the many doesn't cancel your right. War version: civilians have a right not to be targeted — this right stands even when violating it would "help win."

(d) Moral status / respect. You can treat someone in a way that fails to **respect their dignity** even when the *outcome* looks fine. People aren't just containers of good-or-bad results; they're owed respect as people.

Lying to a competent adult "for their own good" can wrong them even if it works out well — you've treated them as a child, not a person who deserves the truth. War version (this becomes Sparrow's killer-robot "respect" argument in Week 11): being killed by a machine that "decided" you should die may fail to respect you as a human, even if the machine targets accurately. The *outcome* (a correct target hit) doesn't, on this view, settle whether you were *respected*. (Note: this is Sparrow's *argument*, not a settled fact — Arkin and others push back, as you'll see in Week 11.)

Notice the pattern: all four tools give the deontologist ways to say "the outcome isn't the whole story." A consequentialist tends to wave them away ("only results matter"). Knowing who reaches for which tool tells you which family is speaking.

5. The big split inside Just War Theory: *ad bellum* vs *in bello*

Most of this unit lives inside **Just War Theory** — the middle-ground view that war is *sometimes* justified, but only if strict conditions are met. (The two extremes: **realism** = morality doesn't apply to war at all; **pacifism** = war is never justified. Just War Theory sits between them.)

Just War Theory splits into two **completely independent** halves. Mixing them up is the classic beginner mistake, so nail it now:

- **Jus ad bellum** = the justice of **going to war** in the first place. *Should this war be fought at all?* (Is there a just cause, like self-defence? Proper authority? Is it a last resort?)
- **Jus in bello** = justice **in how the war is fought**. *Given that fighting is happening, are the methods clean?* The big rule here is **distinction** (also called discrimination): you may only target combatants, never deliberately target civilians.

"Independent" means all four combinations are possible. This is the part to really get:

1. **Just to start, fought cleanly** — a country fights a justified war of self-defence and scrupulously spares civilians. The ideal.
2. **Just to start, fought dirtily** — a country is rightly defending itself against invasion (good *ad bellum*) but then massacres prisoners or guns down a village of civilians (bad *in bello*). Right to fight, war crimes in the fighting. (My Lai — US troops massacring Vietnamese villagers — is the textbook example of an *in bello* atrocity; whether the wider war was just to *start* is a separate, disputed question, which is exactly the point: the two judgments come apart.)
3. **Unjust to start, fought cleanly** — an aggressor invades with no just cause (bad *ad bellum*) but its soldiers carefully obey every rule, sparing all civilians (good *in bello*). The war itself is a crime, but the conduct is restrained.
4. **Unjust to start, fought dirtily** — an aggressor with no just cause that also slaughters civilians. Wrong on both counts.

The lesson: "*Were they right to fight?*" and "*Did they fight right?*" are two separate questions with two separate answers. In the exam, if a question is about *starting* a war, stay in *ad bellum*; if it's about *how* a war is waged (drones, targeting, weapons), stay in *in bello*. Saying which half you're in already shows the examiner you understand the structure.

6. Thirty-second recap: how this bedrock powers the scary modern stuff

Everything in Weeks 9–11 is the Master Key wearing a costume:

- **Nuclear (Hiroshima):** Did saving lives by ending the war early justify deliberately incinerating a city's civilians? That's *consequentialism vs deontology*, plus *intention vs foresight* (were the civilians *targeted*, or merely "collateral"? On most readings of city-bombing they were deliberately hit, so this is exactly what's disputed), inside *jus in bello* (distinction).
- **Drones / targeted killing:** Precise strikes may cut civilian deaths overall (a consequentialist plus), but do they wrongly target people far from any battlefield, or kill bystanders (the *distinction* rule, *rights*)?
- **Autonomous weapons / killer robots (Sparrow, your examiner's own work):** Even if a robot targets *accurately* (good outcome), does being killed by a machine fail to *respect* the victim as a human, and who is *responsible* when it kills wrongly? That's *moral status/respect* and *rights*, not outcomes.

See it? You're not learning four new topics. You're applying *one* toolkit — outcomes vs rules, intention vs foresight, the two halves of Just War — to four new machines.

5 check-yourself questions (with answers)

Q1. Your friend says, "Drone strikes are just wrong, I feel it in my gut." Why does this score zero marks in this unit, and what's the fix? *A: Because it's a feeling, not a mini-argument — descriptive at best, and ethics is normative. The fix is to add a "because": e.g., "Drone strikes are wrong because they routinely kill civilians who pose no threat, violating distinction." Claim plus reason.*

Q2. A general says, "Bomb the enemy's civilians — it'll end the war early and save more lives overall." Name the two families clashing and what each says. *A: Consequentialism vs deontology. Consequentialism: if the tally genuinely saves more lives, do it — outcomes are all that matter. Deontology: no — deliberately killing innocents is wrong in itself and their right not to be murdered trumps the good total, even if the maths works.*

Q3. A pilot bombs a tank factory knowing some neighbours will die, versus bombing a neighbourhood to terrorise the population. Which tool distinguishes these, and is there a moral difference? *A: Intention vs foresight. The first merely foresees civilian deaths as a side-effect; the second intends them as the means. Deontologists say yes, this matters greatly — aiming at innocent death is far worse than foreseeing it.*

Q4. Country A invades its neighbour for no good reason but its soldiers carefully spare every civilian. Which half of Just War Theory does it pass, and which does it fail? *A: It fails *jus ad bellum* (no just cause for going to war — it's an aggressor) but passes *jus in bello* (it fights cleanly, obeying distinction). The two are independent, so a war can be unjust to start yet fought by the rules.*

Q5. A killer robot targets a genuine enemy combatant with perfect accuracy. A consequentialist shrugs — correct target, good outcome. Why might a deontologist still object? *A: Because outcomes aren't the whole story. Sparrow gives two separate objections. (1) The **respect** argument: being killed by a machine that "decided" your death may fail to respect you as a*

human being — no human moral agent is acknowledging the gravity of taking your life. (2) The **responsibility gap**: if the machine kills wrongly, no one is fairly answerable — not the programmer (couldn't predict it), not the commander (didn't decide it), not the machine (you can't blame a machine) — and just war requires that someone can be held responsible. Either way, an accurate outcome doesn't, on Sparrow's view, settle whether the victim was treated as a person is owed. (Worth flagging: Arkin contests this — he argues robots without fear or rage could kill fewer civilians than stressed soldiers, so the debate isn't closed.)