

Deliverable A — The Question Bank + Model Answers

~38 predicted oral questions, tagged Q1 (wks 9-11) / Q2 (wks 1-8) / Q3 (synthesis). Each in POSITION -> REASON -> OBJECTION -> REPLY -> EXAMPLE, speakable in 3-4 min, with likely follow-ups. Practise these ALOUD.

Q1 BANK — The Hard Cases (Weeks 9-11)

Are nuclear weapons mala in se — evil in themselves? [Q1]

- **Position:** Yes — strategic nuclear weapons are evil-in-themselves (wrong by their very nature, not just by misuse), because they cannot in principle satisfy distinction or proportionality.
- **Reason:** "Mala in se" means wrong by nature, not just wrong by its effects on a given day. A strategic nuke incinerates a whole city: it kills masses of non-combatants deliberately — city-busting targets the population itself as a means of coercion — so it violates distinction (only target combatants) by design, every time. And its fallout, radiation and climate effects cross borders and persist for generations, so the harm is uncontainable and can't be weighed proportionately. A weapon that fails the in bello (rules-of-fighting) tests by its nature, every single time, is the textbook case of an evil means — alongside poison or torture. And only a deontologist (some acts wrong in themselves) can say this; a pure consequentialist, who only weighs outcomes, has no category for "evil-in-itself."
- **Objection:** A tiny tactical nuke dropped on an isolated enemy fleet at sea kills only combatants and contains its effects — so it discriminates and is proportionate, which shows nukes can't be mala in se as a category.
- **Reply:** At best that rescues one marginal sub-case, not the strategic city-busting arsenal that actually defines nuclear deterrence — which is still evil-in-itself. And even the "clean" tactical strike carries an irreducible escalation risk: a fleet strike can trigger a full strategic exchange, which poisons the proportionality sum. The unit's reflective exercise also presses us — is radiation death really worse than shrapnel death, or are we just scope-blind? My answer: the wrong isn't the manner of death, it's the in-principle impossibility of confining it to combatants.
- **Example:** Hiroshima — roughly 70,000 people killed instantly, the overwhelming majority civilians, with radiation deaths continuing for years. The strike's entire mechanism is the mass killing of a city's population, which is precisely what distinction forbids.

- **Likely follow-ups:**

- *If a tactical nuke really could hit only combatants at sea with no fallout, would you still call it evil-in-itself?* → No — I'd concede that one use is permissible, but the strategic arsenal that defines nuclear weapons remains mala in se.
 - *Is the wrong in the radiation specifically, or in the scale of indiscriminate killing?* → In the indiscriminate killing — radiation just makes the uncontainability vivid; even a "clean" city-buster would still target non-combatants.
 - *Could a consequentialist agree nukes are evil-in-themselves?* → Not strictly — they can only call them wrong-in-their-effects; "evil-in-itself" is a deontological verdict.
- **If you blank:** A strategic nuke deliberately incinerates a whole city and spreads uncontainable fallout, so it breaks distinction and proportionality by its very nature, every time — that's what makes it mala in se, and only a deontologist can say so.

Can it be ethical to threaten to do something that it would be unethical to actually do? (the deterrence paradox) [Q1]

- **Position:** No — sincere nuclear deterrence is unethical, because it requires genuinely intending to murder millions of innocents, and forming that intention is already a moral wrong, peace or no peace.
- **Reason:** Deterrence (Mutual Assured Destruction) only works if the threat is sincere — the enemy must believe you would really annihilate their cities in retaliation. So a deterrent posture isn't a bluff; it requires a standing, conditional intention to commit mass murder of non-combatants. On a deontological view, what you intend is central to what you're doing — intending to slaughter innocents is already corrupt even if the bombs are never fired. Kavka sharpens this with the wrongful-intentions puzzle (it can be rational to intend to do what it would be wrong to do). So you face a fork — either the wrongful intention taints deterrence, or you say good consequences license a monstrous conditional intention, and that second route is exactly the permissive consequentialism the unit warns against.
- **Objection:** The conditional intention is justified by its consequences — roughly 80 years without great-power war. A wrongful act you never commit, held precisely so you never have to commit it, is surely better than the alternative; intending-in-order-never-to-act isn't the same as murdering.
- **Reply:** That defence stands or falls on the consequences actually obtaining — and the empirical case is shaky. Ward Wilson argues deterrence may not reliably work at all, and Schlosser documents how close we've come to accidental and false-alarm launches ("World War Three by mistake"). So you'd be cultivating a genuine intention to murder millions in exchange for a peace-dividend that may be illusory and a catastrophe-by-mistake that is very real. Strip out the shaky consequences and you're left with the bare wrong of the intention.
- **Example:** The 1983 Petrov incident — a Soviet early-warning system falsely reported incoming US missiles, and standing doctrine was to retaliate. One officer's refusal to pass it up the chain is the only reason the "conditional intention" wasn't executed against tens of millions. It shows the threat was sincere, and how thin the line is between deterrence and accidental mass murder.

- **Likely follow-ups:**

- *If deterrence demonstrably did prevent war, would that make the intention to incinerate cities permissible?* → Only on consequentialism — and I reject that here, because some intentions (to murder innocents) are corrupt whatever they buy.
 - *Is there a morally relevant difference between a bluff and genuine deterrence?* → Yes — a bluff carries no real intention to fire, so it's cleaner morally, but it's unreliable: it only works if the enemy can't tell it's a bluff.
 - *Whose puzzle is this?* → Kavka and Gauthier name it precisely — the wrongful-intentions puzzle, that a wrongful intention may be rational to form, which sharpens rather than dissolves the paradox.
- **If you blank:** Deterrence only works if the threat is sincere, so it requires genuinely intending to murder millions of innocents — and forming that intention is already wrong on a deontological view, while the good consequences that might excuse it (Ward Wilson, Schlosser) are too shaky to bank on.

Was the atomic bombing of Hiroshima ever justifiable? [Q1]

- **Position:** No — Hiroshima was not justifiable, because it deliberately targeted a civilian population (violating distinction outright), and the "it ended the war and saved lives" defence is both forbidden in form and empirically shaky.
- **Reason:** The bombing's mechanism was the intentional killing of a city's inhabitants to coerce a government — that is the deliberate targeting of non-combatants that distinction absolutely forbids. So it can't be rescued by the Doctrine of Double Effect (you may cause foreseen civilian deaths only as a side-effect, never as your means) — here the civilian deaths were the means, not a side-effect. The only serious defence is consequentialist: it shocked Japan into surrender and saved more lives than a land invasion. But that defence both assumes you may murder innocents for the greater good — the move Nagel's absolutism rejects — and assumes the factual claim is even true.
- **Objection:** It was the lesser evil — a land invasion of Japan would have cost far more lives on both sides, so the bombing minimised total suffering; and in total war the combatant/civilian line had already blurred.
- **Reply:** Two problems. First, even if the numbers favoured the bomb, deliberately murdering 70,000 innocents to coerce a third party is exactly the act absolutism says you may never do whatever the arithmetic — otherwise civilian immunity means nothing. Second, the arithmetic itself is contested: Ward Wilson argues it was the Soviet declaration of war and invasion of Manchuria, not the bombs, that drove Japan's surrender — so the lives-saved premise may simply be false, which collapses the only available defence.
- **Example:** Wilson's comparison — many Japanese cities had already been firebombed with comparable death tolls without forcing surrender, yet Japan moved to surrender within days of the Soviet invasion. That suggests it was Soviet entry, not Hiroshima and Nagasaki, that was decisive.

- **Likely follow-ups:**

- *If Wilson is wrong and the bombs really did end the war and save a million lives, would that justify them?* → No — deliberately killing civilians as a means is off the table for an absolutist regardless of the payoff; at most it becomes a tragic choice, not a justified one.
 - *Does it matter that Nagasaki came only three days later?* → Yes — it gave Japan almost no chance to respond to the first bomb, making the second even harder to defend as "necessary."
 - *Why can't Double Effect save it?* → Because DDE permits civilian deaths only as a foreseen side-effect; at Hiroshima the civilian deaths were the very means of coercion, so DDE doesn't apply.
- **If you blank:** Hiroshima deliberately targeted a city's civilians as the means of coercing Japan, which distinction forbids and Double Effect can't excuse — and the only defence, "it saved lives," is both the forbidden consequentialist move and empirically doubted by Ward Wilson, who credits the Soviet invasion.

Does deterrence even work? Does that change the ethics of nuclear weapons? [Q1]

- **Position:** No — I don't think we can be confident deterrence works, and that uncertainty is decisive: it knocks out the only defence nuclear weapons have, so they should be condemned.
- **Reason:** The entire moral case for keeping nuclear arsenals is consequentialist — "they're terrible, but they've kept the great-power peace for 80 years, so the threat is worth it." That argument lives or dies on the factual claim that deterrence actually causes the peace. Ward Wilson ("The Gordian Knot") challenges this directly: the historical record doesn't clearly show nukes preventing wars, and correlation isn't causation — many things kept the Cold War cold. If the deterrent benefit is illusory, then you're holding weapons that are plausibly evil-in-themselves for a payoff you can't bank. Strip out the good consequence and nothing is left to justify the standing intention to incinerate cities.
- **Objection:** The absence of great-power war since 1945 is itself the evidence — no major nuclear power has fought another directly, which is unprecedented, and the most obvious explanation is that the cost of nuclear retaliation made it unthinkable. Demanding a controlled experiment is unfair; the track record speaks for itself.
- **Reply:** A clean 80-year record with no counterfactual is exactly what you'd also see if deterrence did nothing and other factors (economic interdependence, exhaustion, luck) kept the peace — so the record under-determines the conclusion. Worse, Schlosser shows the ledger has a hidden cost: repeated near-misses, false alarms and accidents mean deterrence carries a standing risk of catastrophe-by-mistake. So even on the consequentialist's own terms the sum is shaky — an uncertain benefit against a real, recurring risk of accidental annihilation. That's not a confident "it works."
- **Example:** The Cuban Missile Crisis is read both ways — deterrence advocates say MAD forced both sides to back down, but it's equally a story of how close rational deterrence came to tipping into catastrophe, with decisions made on incomplete information under terror. Pair it with the 1983 Petrov false alarm and the pattern is clear: the system that supposedly keeps us safe repeatedly nearly killed us by accident.

- **Likely follow-ups:**

- *If we proved deterrence works, would you accept nuclear weapons?* → It would strengthen the consequentialist case, but I'd still say the standing intention to murder cities is wrong in itself — so at most I'd soften from "condemn" to "tragic necessity."
 - *Whose argument is the empirical challenge?* → Ward Wilson's "The Gordian Knot" on whether deterrence works and whether the bombs caused Japan's surrender; Schlosser's on accident risk.
 - *Isn't "we can't prove it either way" a reason to keep them, to be safe?* → No — given the catastrophe is mass murder, the burden of proof is on the side that wants to keep the weapons, and uncertainty cuts against them.
- **If you blank:** The whole defence of nukes is "deterrence keeps the peace," but Ward Wilson argues that's empirically doubtful and Schlosser shows the accident risk is real — so the consequentialist payoff is unbankable, which leaves the standing intention to incinerate cities with nothing to justify it.

Is targeted killing more like an act of war or more like extrajudicial execution? [Q1]

- **Position:** It's justified war-killing only when the target is genuinely liable — a combatant or someone responsible for an ongoing unjust threat that can't be stopped by arrest; otherwise it collapses into extrajudicial execution.
- **Reason:** The whole moral licence to kill without trial comes from the war framework, where you may kill combatants because of the threat they pose. If the target is really just a criminal suspect — someone you could arrest — then killing him borrows war's permissions without doing the work war's rules require, which is exactly what "extrajudicial execution" means. So the answer turns entirely on the status of the target — McMahan's liability, meaning moral responsibility for an unjust threat, not the uniform — and not on the weapon or the location.
- **Objection:** But targeted killing happens outside declared armed conflict — there's no battlefield, the person often isn't shooting at anyone right now, and there's no court. By default that looks like law enforcement, where the rule is arrest-not-kill, so calling it "war" is just a convenient label that lets the state execute people it dislikes.
- **Reply:** The "battlefield" isn't a place, it's a relationship of threat. A bomb-maker planning the next attack is a continuing threat even while asleep, just as a soldier is a legitimate target even off the firing line. What does the real work is feasibility of capture: if you can safely arrest him, you must, and killing is wrong; if he's beyond reach in territory no one will police, lethal force can be the only way to stop an ongoing serious threat. So I'm not waving a "war" label to license execution — I'm saying liability plus genuine impossibility of arrest is what moves a case from law-enforcement into defensible killing.
- **Example:** Compare two cases. A confirmed al-Qaeda operational planner in a remote tribal area no government controls, actively directing attacks, who can't be captured — killing him can be defensible war-killing. Versus a man identified only by a "pattern" of phone use and movement in a signature strike, no confirmed identity, no proven threat — that's not war-killing, it's executing a stranger on a hunch.

- **Likely follow-ups:**

- *If capture became possible at the last second, would the strike become impermissible?* → Yes — last resort is a hard condition; if arrest is feasible, lethal force is no longer justified.
 - *Who decides someone is "liable," and does the lack of a court delegitimise it?* → It's a serious worry — Dige's "in principle vs in practice" split bites here: defensible in theory, but dangerous without oversight in practice, so real programmes need scrutiny.
 - *Does it matter that killing is irreversible while a wrongful arrest can be undone?* → Yes — that asymmetry raises the evidential bar; the irreversibility is a reason to prefer capture wherever possible.
- **If you blank:** Targeted killing is defensible war-killing only when the target is genuinely liable (McMahan) and capture is truly impossible; if you could arrest him instead, killing him is extrajudicial execution dressed up as war.

What, if anything, is wrong with signature strikes? [Q1]

- **Position:** Signature strikes are the hardest case to defend, and I think they're usually wrong, because they kill people on the basis of behavioural patterns rather than confirmed identity or proven liability.
- **Reason:** Every defence of targeted killing rests on the target being genuinely liable — responsible for an unjust threat (McMahan). A signature strike deliberately skips that step: you don't know who the person is, only that they fit a profile — military-age male, certain movements, certain associations. That means you are knowingly accepting a real chance of killing an innocent as your target, not merely as a foreseen side-effect — which breaks distinction at its root, not just at the margins.
- **Objection:** In a real war you also kill people whose names you'll never know — you fire on an enemy position without checking IDs. Behaviour is evidence; a man digging by a roadside at night near a known route, carrying what looks like a weapon, is reasonably judged a combatant. Demanding confirmed identity for every strike would make fighting impossible, so signature strikes are just normal battlefield judgement applied through a camera.
- **Reply:** There's a crucial difference. In open battle, behaviour is read in a genuine combat context where the population has largely separated itself from the fighters, so the inference is reasonable. In a drone campaign over a populated area not in declared war, the same "patterns" — men gathering, carrying tools, attending a funeral — are utterly ordinary civilian life. The base rate of innocence is high, so the same signature carries far weaker evidence; acting on weak evidence to kill is precisely the failure of precaution and distinction that JWT forbids. The camera doesn't fix the inference — it can make us overconfident in it.
- **Example:** Reported signature criteria have included military-age males in areas of suspected militant activity, and strikes on gatherings — which have hit weddings and rescuers returning to a strike site (the "double-tap"). Those are people killed for a pattern, not a proven threat — the clearest illustration of why signature strikes break distinction.

- **Likely follow-ups:**

- *Is there any evidential threshold at which targeting-by-pattern becomes acceptable?* → Only if the pattern is overwhelming and the context excludes ordinary civilian life — which in practice almost never holds; confirmed identity should be the default requirement.
 - *How is a signature strike different from police shooting anyone who "looks like" a suspect?* → It isn't, much — and that parallel is exactly why it troubles me; outside genuine combat the law-enforcement intuition (don't kill on a profile) should govern.
 - *If signature strikes on average killed fewer civilians than the alternative, would that make them permissible?* → That's a consequentialist defence; even if true, deliberately accepting unidentified people as targets breaks distinction in the act, which I won't trade away on a body-count.
- **If you blank:** Signature strikes kill by behavioural profile rather than confirmed identity, so you knowingly accept innocents as your targets, not as side-effects — that breaks distinction at the root, which is why they're the hardest form of targeted killing to defend.

Walk me through the café case: fire a cruise missile at terror plotters in a London café, killing 12 innocent bystanders, or let them escape? [Q1]

- **Position:** On those numbers I would not fire — killing 12 confirmed innocents to stop plotters whose attack is uncertain and who might be caught another way fails proportionality and last resort — but the reasoning, not the verdict, is what matters, and it could flip if the threat were certain, catastrophic and capture truly impossible.
- **Reason:** This is a Double Effect and proportionality problem. Under the Doctrine of Double Effect I may cause foreseen civilian deaths as a side-effect of a legitimate strike, but only if they're proportionate to the good achieved and the strike is necessary. Twelve dead innocents is a very heavy cost; the "good" — plotters who might carry out an attack of unknown scale — is speculative; and crucially this is London, a place with police where arrest is a live option. When capture is feasible, lethal force isn't a last resort, so the strike fails on that ground before we even finish the proportionality sum.
- **Objection:** If those plotters succeed they could kill far more than 12 — say hundreds in a city attack. A pure consequentialist says the maths is obvious: 12 now to save 200 later is a bargain, and refusing to fire makes you responsible for the larger death toll you could have prevented.
- **Reply:** Two replies. First, the deontological point: the 12 bystanders aren't liable for anything, so killing them as the means of getting the plotters treats innocents as mere tools, and the intention-versus-foresight line gets strained when the bystander deaths are so massive and certain against a threat so uncertain — at some ratio "side-effect" becomes a fig leaf. Second, even on the consequentialist's own ground the maths isn't clean: the attack is probable-not-certain, its scale is unknown, and a London café is exactly where arrest could stop it with zero innocent deaths — so "fire now" isn't even clearly the harm-minimising option. That's why last resort does the heavy lifting: it's only a genuine dilemma if capture is truly impossible.

- **Example:** Contrast the August 2021 Kabul strike — rushed intelligence, a man wrongly judged a threat, 10 civilians including 7 children killed, and no actual bomb. That's the café case made real: acting on a probable-threat reading and getting the proportionality and the facts catastrophically wrong. It shows why the bar for accepting certain innocent deaths against an uncertain threat must be very high.
- **Likely follow-ups:**
 - *Change the facts — the attack is certain, imminent, will kill thousands, and there's genuinely no way to arrest them. Do you fire?* → Then the calculus shifts and I might fire, because last resort is met and the proportionality sum now clearly favours acting — the verdict tracks the conditions, not a slogan.
 - *Does it matter whether the 12 deaths are your means or a side-effect?* → Yes, hugely — if their deaths are the means, it's deliberate killing of innocents (forbidden); a genuine side-effect can be permissible under Double Effect if proportionate.
 - *If you don't fire and the plotters kill 200, are you responsible?* → I bear some responsibility for the foreseeable outcome, but allowing harm isn't the same as deliberately killing innocents — the doing/allowing distinction means I'm not the murderer of those 200.
- **If you blank:** On those numbers I wouldn't fire — 12 certain innocent deaths against an uncertain threat in a city where arrest is possible fails both proportionality and last resort — but if the threat were certain and catastrophic and capture truly impossible, the verdict could flip; the reasoning is what scores.

Byman says drones work and save civilian lives. Is the precision of drones a moral point in their favour? [Q1]

- **Position:** Yes — precision is a genuine moral point in drones' favour, but only a conditional one: it counts only when the target is genuinely liable and the alternative really would kill more, and it can't license the abuses (signature strikes, casual kill-lists) that drones also enable.
- **Reason:** Byman's case ("Why Drones Work") is serious: a drone can loiter, watch, and strike one liable target with far less collateral harm than an air raid or a ground invasion of the same compound. If the choice is between a precise strike that spares bystanders and a messy raid that doesn't, then refusing the drone is choosing the bloodier option — and minimising harm to innocents is exactly what jus in bello proportionality and distinction demand. So precision genuinely serves the values JWT cares about; it isn't a morally neutral gadget.
- **Objection:** Boyle's reply: this looks good strike-by-strike but is self-defeating overall. Civilian deaths plus the experience of being hunted from the sky radicalise whole populations, creating more enemies than the strikes remove — so "precise" drones are strategically and morally counter-productive once you zoom out from the single target to the campaign.

- **Reply:** Boyle's objection is consequentialist, and that cuts both ways — if blowback genuinely produces net more enemies and deaths, then by the consequentialist's own standard the programme fails proportionality and should stop. But that's an empirical claim that has to be shown case by case, not assumed; in a given case a precise strike on a serious threat that spares the civilians a ground assault would have killed can still come out ahead. The real lesson is that precision is a tool, not a permission: it makes harm-minimising killing possible, but the same technology tempts states to drop the safeguards — killing on suspicion, at zero risk — and that's where it goes wrong, not in the precision itself.
- **Example:** Set Byman's ideal against the August 2021 Kabul strike: rushed intelligence, a man wrongly judged a threat, 10 civilians including 7 children dead, no bomb. The drone was "precise" in the engineering sense — it hit exactly what it aimed at — but precision is worthless when the targeting judgement is wrong. That shows precision is a virtue only downstream of getting liability and the facts right.
- **Likely follow-ups:**
 - *So are you pro-drone or anti-drone?* → Conditionally pro — for confirmed liable targets where capture is impossible and the alternative is bloodier; anti for signature strikes and zero-scrutiny kill-lists.
 - *Is Boyle's blowback point decisive?* → Not on its own — it's an empirical bet that must be demonstrated; if it holds it defeats drones on consequentialist grounds, but it can't be assumed for every case.
 - *Does precision ever make us kill more, not less?* → Yes — that's the worry: low cost and zero risk lower the threshold, so states strike in cases they'd otherwise have left alone.
- **If you blank:** Byman is right that a precise drone strike on a liable target can spare more innocents than a raid, so precision is a real moral plus — but only when the target is genuinely liable and capture is impossible; Boyle's blowback worry and the temptation toward signature strikes show precision is a tool, not a blanket permission.

Is there anything wrong with riskless, asymmetric killing — fighting an enemy who can't fight back? (Sparrow's worry) [Q1]

- **Position:** No — risklessness by itself isn't what makes a killing wrong; the discomfort Sparrow points to is real but it actually attaches to whether the target is genuinely liable, not to the risk gap.
- **Reason:** Force-protection is normally a virtue, not a vice. We don't tell snipers to step into the open, and a just combatant has no duty to share his enemy's danger — keeping your own people safe while stopping an unjust threat is good, not suspect. If the target is genuinely liable (morally responsible for an unjust threat, in McMahan's sense), then killing him at low risk to yourself is permissible defence done efficiently. An absurd "fairness" demand — take off your body armour to make it even — would only get more of your own side killed for no moral gain. So asymmetry per se doesn't taint the act.

- **Objection:** Sparrow's worry: war arguably involves mutual jeopardy — soldiers may kill each other because each poses a danger to the other, and that reciprocal risk is part of what makes battlefield killing different from murder. A drone operator thousands of miles away faces zero danger, so this isn't combat at all; it's one-sided hunting, and pure hunting of humans looks morally troubling even when the quarry is "liable."
- **Reply:** I take the worry seriously, but I think it mislocates the wrong. What licenses killing in war isn't that both sides are in danger — it's that the target is liable, a threat that may justly be stopped. Mutual jeopardy is usually present, but it's a feature of typical battle, not the thing that does the moral work; if it were, then killing a sleeping enemy soldier would be murder, which it plainly isn't — he's still liable, still part of the threat. So the genuine unease about drone killing should redirect to the real questions — is this person actually liable? is the strike proportionate and a last resort? — rather than to the risk gap, which on its own is just successful self-protection.
- **Example:** Compare a sniper who kills an enemy officer from a mile away, unseen and unthreatened, with a drone operator who does the same from a control room. Almost no one calls the sniper a murderer — yet his situation is also riskless and asymmetric. The only morally relevant difference between the two cases is whether the target is genuinely liable and the strike necessary, which is exactly where the scrutiny belongs.
- **Likely follow-ups:**
 - *So you think Sparrow is wrong?* → On this specific point, yes — but he's right that drone killing is troubling; I just think the trouble is liability and oversight, not the absence of risk to the operator.
 - *Doesn't zero risk lower the threshold for going to war?* → That's a real and separate worry — cheap, riskless killing can make states reach for force too readily — but that's an *ad bellum* problem, not proof the *in bello* killing itself is wrong.
 - *Is there any case where asymmetry itself is the wrong?* → Perhaps where it becomes pure hunting of the non-liable — but then the wrong is killing the non-liable, which asymmetry only makes easier, not the asymmetry as such.
- **If you blank:** Sparrow worries that riskless drone killing isn't real combat because war involves mutual jeopardy — but I think force-protection is a virtue and the real moral work is done by liability, not risk; the genuine unease should redirect to "is this target liable and the strike necessary?", not to the risk gap itself.

What is an autonomous weapon, and does it really matter morally whether a human or a machine pulls the trigger? [Q1]

- **Position:** Yes — it matters morally; an autonomous weapon should not be permitted to kill, because the trigger-puller is exactly where responsibility and respect live.

- **Reason:** First the definition: an autonomous weapon system selects and engages targets without meaningful human control — it picks who dies, not just how. That contrasts with a drone, where a machine flies but a human still pulls the trigger and so still chooses each death. Why the trigger matters: the rules of fighting justly (*jus in bello*) buy us the licence to kill at the price of accountability — when a killing turns out wrongful, someone must be justly answerable. The human at the trigger is that someone, and is also the agent who, in deciding, registers that ending a life is grave. Remove the human and you remove both at once — this is Sparrow's responsibility gap plus his respect argument. So "human vs machine" isn't a technical detail, it's the whole moral hinge.
- **Objection:** Arkin: the trigger-puller is the problem, not the solution. Humans pull triggers in fear, rage and revenge — that's what produced My Lai. A machine with an "ethical governor" pulls the trigger only when the laws of war are satisfied, so handing it the trigger could kill fewer innocents. On that view the human is the unreliable element.
- **Reply:** Grant that machines might be more consistent — it still doesn't answer the question. A more reliable trigger-puller is still, when it errs, a killing no one answers for and no one weighed. Arkin is scoring on outcomes; the worry is about the act and its accountability. Fewer wrongful deaths is not the same as wrongful deaths someone can be held to. So even a "better" machine trigger leaves both gaps — responsibility and respect — wide open.
- **Example:** The August 2021 Kabul drone strike killed ten civilians, including seven children. Today we can name and in principle hold responsible the humans who approved it — there's an agent the wrong attaches to, an inquiry, an apology owed. Now imagine an autonomous weapon had selected and struck that target on its own: the same ten dead, but no human chose them, the programmer couldn't have predicted this strike, and the machine can't be blamed. Same bodies, but the accountability has evaporated — that's the difference the trigger makes.
- **Likely follow-ups:**
 - *If a human is "on the loop" who can veto but rarely does, is that meaningful control or a rubber stamp?* → If he can't realistically scrutinise each kill, it's a rubber stamp — meaningful control requires genuine, exercisable judgement, not a nominal veto.
 - *Does your view ban defensive systems too, like a ship's gun that auto-fires at incoming missiles?* → Less so — point-defence against incoming munitions targets objects, not chosen human victims, so the respect and responsibility worries are far weaker there.
 - *If the soldier is far less reliable than the machine, isn't keeping the human in charge knowingly choosing more dead civilians?* → That's Arkin's challenge, and I accept the cost — because severing accountability and respect is a wrong in the act, not a price I can pay down with a lower body-count.
- **If you blank:** An autonomous weapon selects and engages targets with no human pulling the trigger; the trigger is where accountability and respect live, so removing the human opens Sparrow's responsibility gap and his respect problem at once — which is why it matters morally, not just technically.

Explain Sparrow's responsibility-gap argument against autonomous weapons. Is it convincing? [Q1]

- **Position:** It's largely convincing — the responsibility gap is a genuine, distinctively deontological reason to ban autonomous weapons, not just a complaint about bad outcomes.
- **Reason:** The argument runs in three steps. (1) Jus in bello presupposes that for any wrongful killing, someone can be justly held responsible — accountability is the moral price of the licence to kill. (2) For a genuinely autonomous weapon, no candidate qualifies. The programmer couldn't predict the specific act — that unpredictability is precisely what makes it autonomous rather than a tool, so blaming her is like blaming a parent for an adult child's crime. The commander didn't make this decision; punishing her for a choice she neither made nor could foresee is unjust, like jailing an officer for a soldier she had no way to control. The machine can't be blamed or punished — it isn't a moral agent, can't understand the wrong, can't suffer, has nothing to be sorry for; "punishing" it is theatre. (3) So you have a wrongful killing with no just bearer of responsibility — a gap — and creating that gap deliberately makes the killing illegitimate.
- **Objection:** Strict liability closes the gap — just decree in advance that the commander (or the manufacturer) owns every outcome, like a dog owner is liable for the dog's bite. Problem solved: there's now always someone answerable.
- **Reply:** That assigns blame without responsibility, which is itself unjust. We hold people responsible for what they could control or foresee; punishing a commander for a death she couldn't predict isn't accountability, it's a scapegoat with a name tag. Strict liability might satisfy a lawyer's need for a defendant, but Sparrow's point is moral, not legal — it doesn't make her genuinely answerable, it just papers over the gap with an injustice. And note the deeper move: making the weapon more autonomous — smarter, more independent — widens the gap, because it pushes the act further beyond anyone's foresight. The better the robot, the worse the problem.
- **Example:** Compare a landmine and an autonomous hunter-drone. A landmine is dumb — whoever laid it chose the kill-zone, so responsibility traces cleanly back to a human decision. An autonomous drone that roams, identifies a target by its own learned criteria, and strikes is different: no human chose this victim, and no human could have predicted the machine would read this person as a threat. The landmine has an author; the truly autonomous strike has none. That missing author is the gap.
- **Likely follow-ups:**
 - *If future law simply declared the deploying officer liable for everything, would that satisfy you?* → No — it just renames the problem; legal liability without genuine foresight or control isn't moral responsibility, it's a scapegoat.
 - *Is the gap really new — when an artillery shell goes astray and kills civilians, who's justly responsible?* → There the human chose to fire into that situation and accepted the known risk, so responsibility traces back; with a truly autonomous weapon no human chose this victim at all.
 - *Does the argument prove too much — would it ban any complex system like an air-defence network?* → No — those still act on human-set parameters against incoming threats; the gap bites specifically where the machine itself selects human targets by its own learned criteria.

- **If you blank:** Jus in bello requires that someone be justly answerable for a wrongful killing, but with a truly autonomous weapon no one qualifies — not the programmer (couldn't foresee it), not the commander (didn't choose it), not the machine (no moral agent) — so deploying it creates an accountability gap, and that makes the killing illegitimate.

Arkin argues robots could be MORE humane than soldiers. Should we build autonomous weapons if they would kill fewer civilians? [Q1]

- **Position:** No — even if autonomous weapons killed fewer civilians, we should not build them; a lower body count does not dissolve the responsibility gap or restore the respect owed to the people they kill.
- **Reason:** This question is the unit's master key — consequentialism versus deontology — in a robot costume. Arkin argues from outcomes: machines have no fear, anger, revenge or self-preservation (the emotions behind My Lai), an "ethical governor" could obey the laws of war more consistently than a stressed soldier, so if AWS kill fewer innocents, refusing them is the choice that kills more — banning them would be the immoral act. You must state that fairly. But Sparrow's two objections are deontological — about the kind of act, not its tally. (1) Fewer wrongful killings still leaves each remaining wrongful death with no one justly answerable. (2) Being efficiently killed by something that cannot grasp that your death matters is its own wrong, no matter the statistics. Good consequences can't pay a deontological debt, because the debt isn't about consequences.
- **Objection:** That sounds like fetishising procedure over lives. If you genuinely could save a thousand civilians a year by deploying ethical-governor robots, refusing them on grounds of "accountability" and "respect" means choosing those thousand deaths for a philosophical scruple. Surely the dead would rather be alive-and-disrespected.
- **Reply:** Two replies. First, the empirical premise is shaky — the hard part of the laws of war is proportionality, weighing military advantage against civilian harm, which is a judgement, not an object-recognition task; a machine that can spot a rifle still can't weigh whether this target is worth those deaths, so "fewer civilians" is a promise, not a fact. Second, even granting it: I'm not refusing to save lives, I'm refusing a specific method of killing that severs accountability and respect — the same reason we ban torture even where torture might extract life-saving information. We do think some ways of getting good outcomes are off-limits, and AWS-killing is one of them.
- **Example:** Take My Lai, Arkin's own touchstone. His point: a robot wouldn't have snapped and massacred a village out of rage. True. But the reason My Lai haunts us isn't only the count of the dead — it's that humans did it, were named, were tried, and a society had to reckon with the wrong. An autonomous massacre with no perpetrator to try, no one to be sorry, would be cleaner in the ledger and morally emptier — the killing of people that, to the killer, never registered as mattering. Fewer dead, but no one home to answer for them.

- **Likely follow-ups:**

- *Where exactly is your line — all autonomous lethal force, only against humans, or only offensive use?* → Against humans, where selecting the target is a moral judgement; point-defence against incoming munitions is far less troubling.
 - *If a machine really could make proportionality judgements as well as a good officer, would your objection collapse to just respect?* → Largely yes — the gap worry would ease, but the respect argument still bites independently, which is why Sparrow keeps the two separate.
 - *Doesn't the respect argument also condemn high-altitude bombing, where the killer never sees the victim?* → It strains there, but a bomber pilot still chose to drop on that target and weighed it; with AWS no human ever weighs this death at all — that's the difference.
- **If you blank:** This is consequentialism vs deontology in a robot costume — Arkin says fewer civilian deaths makes AWS the humane choice, but Sparrow's objections are about the act: a lower body count still leaves wrongful deaths no one answers for and people killed by something that never registered their death as mattering, so good outcomes can't pay that debt.

Can a machine make the judgements jus in bello requires — distinction and proportionality? [Q1]

- **Position:** No — distinction is hard and proportionality is the real wall: it's a moral judgement, not a recognition task, so a machine can't make the calls jus in bello requires, which is an independent reason to keep humans in control.
- **Reason:** Distinction (only target combatants, never deliberately the innocent) is partly a perception problem — telling a fighter from a farmer — and a machine might get better at the easy version, spotting a rifle or a uniform. But liability isn't a shape; McMahan's point is that what makes you a legitimate target is moral responsibility for an unjust threat, and a camera can't read responsibility. Proportionality is worse still: it asks whether the military advantage of a strike outweighs the foreseeable civilian harm — that's weighing incommensurable goods, a value judgement embedded in context, not a number you compute from pixels. A system that can identify a target still can't judge whether this target is worth those deaths. So the genuinely moral parts of jus in bello are exactly the parts code can't do.
- **Objection:** Arkin's reply: don't overstate what humans do either. A frightened, exhausted, vengeful soldier makes those judgements badly — that's how My Lai happens. A machine with an "ethical governor," applying the rules consistently and without panic, could make distinction and even rough proportionality calls more reliably than a stressed human, so "machines can't judge" sets the bar in the wrong place.
- **Reply:** Two things. First, consistency isn't the same as correctness on a value judgement — a machine can apply a rule flawlessly and still apply the wrong rule, because proportionality has no fixed formula to encode; the weighting itself is the contested moral act, so reliably executing a bad weighting is worse, not better. Second, even if a machine matched a good officer on the numbers, that only neutralises the capability worry — it leaves Sparrow's responsibility and respect arguments untouched, because those are about who answers for the death and who weighed that it mattered, not about getting the arithmetic right. So "the machine can judge" wouldn't rescue AWS even if it were true.

- **Example:** Picture two armed men beside a building. One is an enemy combatant; the other is a civilian guard who'd lay down his weapon if asked — Sparrow's own "hands up who wants to die?" intuition is that a human can offer surrender and read the moral situation. A machine sees two armed shapes and a rule; it can't extend the offer, can't read the difference in responsibility, and can't weigh whether the building's military value justifies killing both. That gap between pattern-matching and moral judgement is the whole problem.
- **Likely follow-ups:**
 - *If AI got good enough to match human proportionality calls, would you allow AWS?* → No — it would remove the capability objection but leave the responsibility gap and the respect argument standing, and those are decisive.
 - *Isn't distinction just object recognition, which machines are great at?* → The easy version, maybe — but liability is moral responsibility, not a visible feature, so confident object recognition can still target the wrong people.
 - *Whose distinction is "liability"?* → McMahan's — legitimate targeting tracks moral responsibility for an unjust threat, not the uniform, which is exactly what code can't see.
- **If you blank:** Distinction's easy version is recognition, which machines might do — but liability (McMahan) is moral responsibility, not a shape, and proportionality is a value judgement weighing military gain against civilian harm, not a calculation; those moral calls are what code can't make, which is an independent reason to keep humans in control.

Do drones and autonomous weapons lower the threshold of war — make states too willing to fight? [Q1]

- **Position:** Yes — by making war cheap, riskless and politically invisible, drones and autonomous weapons lower the threshold for resorting to force, and that's a serious moral cost that sits on the ad bellum side of the ledger, separate from any single strike.
- **Reason:** Just war theory splits into two independent halves: jus ad bellum (the justice of going to war) and jus in bello (justice in fighting it). Most drone debate is about in bello — is this strike discriminate and proportionate? But there's a distinct ad bellum worry: war is supposed to be a last resort, and one of the brakes on it has always been cost — your own soldiers come home in coffins, and the public notices. Remove the risk to your own side and remove the domestic political price, and you weaken last resort itself. A tool that lets a state kill abroad with no body bags and little scrutiny makes reaching for force too easy, so even perfectly clean individual strikes can add up to a state that fights when it shouldn't.
- **Objection:** This proves too much. By the same logic any technology that reduces your own casualties — body armour, stand-off missiles, a stronger army — "lowers the threshold," yet we don't say armour is immoral. Reducing harm to your own people is good; if each individual strike is just (liable target, proportionate, last resort), the aggregate worry is speculative and shouldn't override real, defensible strikes.

- **Reply:** The difference is degree becoming kind. Body armour still sends soldiers into danger and home in coffins; drones and AWS can remove the human cost and the public visibility almost entirely, and it's precisely that political cost which historically forced governments to treat war as a genuine last resort. So this isn't a complaint about any one strike — each may pass in bello — it's that the practice erodes the *ad bellum* brake that keeps states from fighting too readily. And the worry is empirical but not idle: a drone programme that quietly runs for years across several countries, with no declared war and little debate, is exactly what "threshold lowered" looks like in practice.
 - **Example:** The US drone campaign across Pakistan, Yemen and Somalia — covered by Coll's "The Unblinking Stare" — ran for years as a persistent, low-visibility killing programme outside any declared war, with minimal domestic political cost because no American was at risk. That's the threshold-lowering worry made concrete: sustained lethal force that a casualty-bearing, publicly-debated war would likely never have authorised.
 - **Likely follow-ups:**
 - *Is this an in bello or ad bellum objection?* → *Ad bellum* — it's about going to war too readily (last resort), not about how individual strikes are fought, so it's logically separate from the precision debate.
 - *Doesn't reducing your own casualties count as a good?* → Yes, in bello it's a virtue — but the same feature weakens the *ad bellum* brake of cost and visibility, so one technology can be good on one axis and corrosive on the other.
 - *How does this connect to autonomous weapons?* → Even more sharply — AWS remove not just risk but human decision-makers, so the political and moral friction that slows a state down nearly vanishes.
 - **If you blank:** This is an *ad bellum* worry, not an *in bello* one: war's cost — your own dead and public scrutiny — has always helped enforce last resort, and drones plus AWS remove that cost, so they make states too willing to fight even when each individual strike looks clean.
-

Q2 BANK — Foundations & Just War Theory (Weeks 1-8)

What is the difference between *jus ad bellum* and *jus in bello*, and why does the distinction matter? [Q2]

- **Position:** The two are genuinely independent questions, and keeping them separate is essential to thinking clearly about war.
- **Reason:** *Jus ad bellum* asks whether it was just to *go to war* (just cause, legitimate authority, right intention, *ad bellum* proportionality, last resort, prospect of success); *jus in bello* asks whether the war is being *fought* justly (*in bello* proportionality, no evil means, necessity). Because they are independent, a war can be just to start yet fought with atrocities, or unjust to start yet fought "cleanly." If you fused them you would wrongly conclude that a side with a just cause may do anything, or that a side with an unjust cause may do nothing.

- **Objection:** A skeptic might say the line collapses in practice — surely an unjust war taints everything done in it, so the conduct rules don't really float free.
- **Reply:** That confuses moral evaluation of the war with the rules that protect people *inside* it. Even in an unjust war, deliberately gassing civilians is a distinct, additional wrong, and a soldier on the just side still murders if he tortures prisoners. The protections only work if they apply regardless of which side you're on.
- **Example:** The 2003 Iraq invasion: many argue it failed *jus ad bellum* (no genuine just cause), yet that's a *separate* question from how it was fought — the Abu Ghraib prisoner abuse was an *in-bello* wrong (mistreating people under your control) stacked *on top of* any ad-bellum wrong.
- **Likely follow-ups:**
 - *Which half does "last resort" belong to?* *Jus ad bellum* — it's about whether to start.
 - *Which half is "distinction" in?* *Jus in bello* — it governs how you fight.
 - *Can the just side commit war crimes?* Yes; ad bellum justice gives no licence in bello.
- **If you blank:** Ad bellum is about *starting* a war, in bello is about *how you fight it*, and they're judged separately.

Why do most philosophers in this unit reject ethical relativism? [Q2]

- **Position:** Relativism has to be set aside, because if it were true we couldn't actually condemn even the clearest atrocity, and that's a decisive reason to reject it.
- **Reason:** Relativism says moral claims are only ever true "for a culture" or "for an individual" — mere taste. But the whole unit assumes some answers are genuinely better than others. If relativism held, "the My Lai massacre was wrong" would mean only "I personally dislike it," not that anyone *ought* not have done it. That guts ethics of its point, which is to give reasons that bind across people.
- **Objection:** Different cultures plainly disagree about morality, and who are we to impose our standards — isn't tolerance itself a reason to be relativist?
- **Reply:** Disagreement doesn't prove there's no fact, any more than disagreement about the cause of a disease means there's no fact. And "you must tolerate" is itself a non-relative moral claim — so relativism can't even state its own tolerance ideal without contradiction.
- **Example:** My Lai (1968), where US soldiers killed hundreds of unarmed Vietnamese villagers. We want to say this was *wrong*, full stop — not just unfashionable. Only a non-relativist can say that.
- **Likely follow-ups:**
 - *Isn't this just Western values?* The argument is structural, not cultural — relativism can't condemn massacre on anyone's terms.
 - *Who said this matters here?* It's the unit's ground rule (the MacKinnon reading sets it up).
 - *Does rejecting relativism mean certainty?* No — just that some answers are better, reachable by argument.
- **If you blank:** If relativism were true we couldn't call any massacre wrong, only say we dislike it — so we reject it.

What is the strongest version of pacifism, and what is the best objection to it? [Q2]

- **Position:** The strongest pacifism is the *contingent* (consequentialist) kind, but it still fails, because some defensive wars really do prevent greater evils.
- **Reason:** It helps to split pacifism into two. *Absolute* pacifism says killing is wrong in itself, always — a deontological line (wrong whatever the result), e.g. Tolstoy's renunciation of all violence. *Contingent* pacifism says war's costs *always in fact* outweigh its gains — an empirical, consequentialist bet (judged by outcomes). The contingent version is stronger because it doesn't have to claim you may not even kill to stop a murderer; it just claims, soberly, that war never pays. But that claim is falsified by cases where fighting clearly averted something far worse.
- **Objection:** The pacifist replies that we systematically underrate war's hidden, long-run costs — the trauma, the precedent, the next war it seeds — so even "good" wars are net negative once you count honestly.
- **Reply:** Sometimes true, but it can't be true *by definition* of every war. A war that stops a genocide in progress can prevent more death than it causes; to insist otherwise is to prejudge the arithmetic rather than examine it.
- **Example:** Allied resistance to Nazi expansion: hard to argue that *not* fighting would have produced fewer deaths or less suffering than fighting did.
- **Likely follow-ups:**
 - *Which pacifism is harder to refute?* Contingent — it's a factual claim, not an absolute prohibition.
 - *Who's the classic pacifist?* Tolstoy, for the absolute, patriotism-renouncing version.
 - *Is pacifism a kind of consequentialism?* The contingent form is; the absolute form is deontological.
- **If you blank:** The best pacifism says war never pays off, but some defensive wars demonstrably prevent worse — so it fails.

Anscombe calls pacifism "pernicious." What does she mean, and is she right? [Q2]

- **Position:** Anscombe is largely right: pacifism is dangerous because it erases the line between killing combatants and murdering the innocent.
- **Reason:** In "War and Murder," Anscombe — a just-war Catholic, not a pacifist — argues that pacifism lumps *all* killing together as equally forbidden. Once you say killing an enemy soldier is just as wrong as killing a baby, you've destroyed the most important distinction in the ethics of war. And in practice, people who can't make that distinction, but who do end up fighting anyway, lose the resources to say "this far and no further" — so pacifism perversely loosens the restraints on the very violence it opposes.
- **Objection:** That's unfair — the pacifist isn't blurring the line, they're refusing to cross *any* line; surely a blanket "no killing" is *more* restrictive, not less.
- **Reply:** As a personal stance, perhaps. But as public morality it's self-defeating: if all killing is murder, then a soldier who has already "fallen" by fighting has no reason to spare civilians — he's damned either way. The combatant/innocent distinction is what *does* the restraining, and pacifism throws it out.

- **Example:** A soldier taught only "all killing is murder" has no principled reason, once in combat, not to bomb a village — whereas distinction tells him soldiers are targets and civilians are not.
- **Likely follow-ups:**
 - *Is Anscombe a pacifist?* No — she defends just war and the Doctrine of Double Effect (foreseen vs intended harm).
 - *What does she want to protect?* The distinction between combatants and the innocent.
 - *Does she think war is fine?* No — she's fierce about not targeting the innocent, which is why she attacks pacifism *and* terror-bombing.
- **If you blank:** Anscombe says pacifism is pernicious because it erases the combatant/innocent line that actually restrains violence.

Is the realist right that morality is simply irrelevant to war? [Q2]

- **Position:** No — the realist is wrong, and the giveaway is that realists keep smuggling moral judgments back in.
- **Reason:** Realism (Morgenthau, Bacevich) says ethics has no place in international affairs: states pursue power and survival in an anarchic world (no overarching ruler), and a state that fights morally while its enemies don't will simply lose. But there's a crucial split. *Descriptive* realism merely says leaders *do* act on interest — that may be partly true but proves nothing about what they *ought* to do. *Prescriptive* realism says they *should* ignore morality — and that is itself a moral "ought," so it refutes the claim that morality doesn't apply.
- **Objection:** The realist replies that this is naive: in a world of rival states, a "moral" state is a dead state, so survival simply trumps ethics — that's not a moral claim, it's prudence.
- **Reply:** As Orend argues, "you ought to ensure your state survives" is still a value claim, and we routinely *do* judge states morally — Nuremberg held leaders responsible for aggression. Realists also quietly invoke fairness, treaties, and outrage at atrocities, which they couldn't do if morality truly didn't apply.
- **Example:** Nuremberg: the world prosecuted Nazi leaders for waging aggressive war — incoherent if morality genuinely had no purchase on statecraft.
- **Likely follow-ups:**
 - *Descriptive vs prescriptive realism?* One describes how states act; the other says they ought to ignore ethics — and "ought" is moral.
 - *Who rebuts realism?* Orend, from a just-war standpoint.
 - *Strongest realist point?* That a purely moral state may be exploited — but that's prudence, not proof ethics is irrelevant.
- **If you blank:** Realists say morality doesn't apply to war, but "states *should* ignore ethics" is itself a moral claim — so it self-defeats.

Is national self-defence morally analogous to individual self-defence? [Q2]

- **Position:** No — Rodin is right that the analogy breaks down, so the standard self-defence justification for war is shakier than Walzer suggests.

- **Reason:** Walzer's "legalist paradigm" (treating states like persons under law) treats a state like a person: aggression against its sovereignty and territory is like an attack on a person's body, and just as you may kill to defend your life, a state may go to war to defend itself. Rodin's challenge is that the analogy fails at the edges. In *bloodless conquest* — where an aggressor takes over without killing anyone — there's no threat to individual lives, yet the standard view still permits a bloody defensive war. So national defence must be protecting something *other* than lives — political institutions and self-determination — and it's far less obvious you may kill on a large scale to defend *those*.
- **Objection:** But losing your state really does threaten lives and freedoms over time — conquest enables future oppression, so defending the state *is* defending persons, just downstream.
- **Reply:** Sometimes, yes — but then the justification rides on those downstream harms, not on a clean person-state analogy. And where conquest is genuinely benign, the analogy gives no warrant for killing, which is exactly Rodin's point.
- **Example:** Imagine a powerful neighbour annexing a small state without firing a shot, leaving daily life untouched. Walzer's paradigm still licenses a lethal war of "defence"; Rodin asks what, exactly, justifies the killing.
- **Likely follow-ups:**
 - *Whose paradigm is being challenged?* Walzer's legalist/self-defence paradigm.
 - *What's Rodin's label?* The reductivist challenge — reduce war's ethics to individual rights and the gap shows.
 - *Which jus ad bellum criterion is in play?* Just cause — what counts as a wrong worth killing over.
- **If you blank:** Rodin says bloodless conquest threatens no lives yet we still permit lethal defence — so war isn't cleanly like personal self-defence.

How should proportionality be judged in deciding whether to go to war? [Q2]

- **Position:** Proportionality is indispensable but genuinely hard, because it requires weighing goods of different kinds — and Hurka is right that *which* goods count is the real battleground.
- **Reason:** Ad bellum proportionality (a jus ad bellum criterion — about whether to start) says a war's expected good must outweigh its expected harm. The trouble, as Hurka stresses, is the weighing: do you count only the just cause itself, or also "relevant" goods like deterring future aggressors, and may you count bad effects on the *enemy's* civilians and soldiers? Get the ledger wrong and proportionality either licenses almost anything (count every speculative good) or forbids everything (count every speculative harm).
- **Objection:** This vagueness makes proportionality useless — if reasonable people fill the ledger differently, it can justify any war the leader already wanted.
- **Reply:** Vagueness isn't uselessness. Proportionality still rules out the clear cases — a vast war over a trivial slight fails on any sane accounting — and it forces leaders to *show their working*, which is exactly what a moral check should do.
- **Example:** Responding to a minor border skirmish with a full-scale invasion costing thousands of lives plainly fails ad bellum proportionality, however you weight the goods.

- **Likely follow-ups:**

- *Is this ad bellum or in bello proportionality?* This is ad bellum — the whole war's costs vs gains; in bello proportionality weighs a single strike.
 - *Who wrote on this?* Hurka, on which goods may be counted.
 - *Does it count enemy lives?* Most just-war theorists say yes — enemy civilian harm goes in the ledger.
- **If you blank:** Proportionality weighs a war's likely good against its likely harm; the hard part, per Hurka, is deciding which goods even count.

Does the "Responsibility to Protect" justify overriding a state's sovereignty? [Q2]

- **Position:** Yes in principle — sovereignty shouldn't shield mass atrocity — but the doctrine is so abusable that I'd apply it only with great caution.
- **Reason:** The classic view treats sovereignty (a state's right to run its own affairs) as near-absolute: states run their own affairs free of outside force. Annan reframes it as "sovereignty as responsibility" — a state forfeits its protection from intervention when it grossly fails to protect its own people, grounding the Responsibility to Protect (R2P). That's the right instinct: sovereignty exists to serve people, not to license a government to slaughter them. So R2P can, in principle, justify humanitarian war.
- **Objection:** Coady and others warn that R2P is a standing pretext: powerful states invoke "humanitarian" motives to pursue strategic interests, and Krasner calls sovereignty "organised hypocrisy" — invoked or ignored as it suits the strong.
- **Reply:** That's a reason for strict tests — genuine atrocity, right intention, proportionality, real prospect of success — not for abandoning R2P entirely. Refusing all intervention also has victims; the answer is rigour, not a blanket rule either way.
- **Example:** NATO in Libya (2011): launched under an R2P banner, but Kuperman's reassessment argues it overstated the threat and worsened outcomes — a cautionary case for how the doctrine gets abused.
- **Likely follow-ups:**
 - *Who coined "sovereignty as responsibility"?* Annan, behind R2P.
 - *What's Krasner's phrase?* Sovereignty as "organised hypocrisy."
 - *Is your nationality deserved?* Arguably not — which weakens treating borders as morally decisive against rescue.
- **If you blank:** Sovereignty serves people, so it can be overridden to stop atrocity — but R2P is so abusable it needs strict tests.

What makes a person a legitimate target in war — being a threat, being responsible, or wearing a uniform? [Q2]

- **Position:** Moral responsibility for an unjust threat is the right criterion — McMahan's liability view — not the uniform.

- **Reason:** This is the heart of *distinction* (only target combatants): never deliberately the non-combatant. But *why* are civilians immune? Three candidate grounds: they are *innocent* (morally blameless), *harmless* (pose no threat), or *not responsible* (haven't contributed to an unjust threat). McMahan argues the deep ground is liability: you become a legitimate target by being *morally responsible for an unjust threat* — which is what "innocent" really tracks, the *not-responsible* sense, not your job title. The uniform is just a rough proxy; what does the moral work is responsibility for the threat.
- **Objection:** A responsibility-based account, pushed hard, threatens civilian immunity: if responsibility is the test, then voters, taxpayers, and munitions workers who *support* an unjust war seem partly responsible — so they lose immunity, which looks monstrous.
- **Reply:** Sparrow's reply ("Hands up who wants to die?") defends a *robust* immunity: responsibility for war comes in degrees, and the contribution of an ordinary civilian is so diffuse and small that it can't ground a *right to kill* them. So responsibility explains liability without dissolving civilian protection — we keep a bright line around non-combatants for good moral and practical reasons.
- **Example:** The Brereton report's alleged killings of Afghan non-combatants: whatever their distant "support" for anything, they posed no threat and bore no responsibility for an unjust threat — squarely protected.
- **Likely follow-ups:**
 - *Whose liability view is this?* McMahan's — responsibility, not uniform.
 - *Who defends robust immunity?* Sparrow, replying to responsibility-based accounts (incl. Primoratz) that risk eroding it.
 - *Is "harmless" enough on its own?* Not quite — a sleeping unjust-side soldier is momentarily harmless but still liable through his responsibility for the threat; responsibility captures the cases better.
- **If you blank:** What makes you targetable is moral responsibility for an unjust threat (McMahan), not your uniform or mere harmfulness.

Does the Doctrine of Double Effect let armies kill civilians "through the back door"? [Q2]

- **Position:** No — properly applied, Double Effect is a genuine moral limit, not a loophole, because it forbids *intending* civilian deaths.
- **Reason:** The Doctrine of Double Effect (DDE), defended by Anscombe and Nagel, says you may cause civilian deaths you *foresee* but don't *intend*, if they're a side-effect of a legitimate strike and proportionate to the military gain. The whole point turns on *intention vs foresight*: aiming at civilians' deaths — as your end or your means — is always forbidden; accepting them as a regretted side-effect can be permissible. So far from a loophole, DDE is what lets just-war theory allow *some* collateral harm without abandoning the absolute ban on deliberately killing the innocent.
- **Objection:** The line between "intending" and "merely foreseeing" is too slippery — a commander can always *redescribe* his aim so the dead civilians become an unintended side-effect, gaming the doctrine.

- **Reply:** The proportionality condition does real work here: even foreseen deaths must be outweighed by the military advantage, so you can't hide a massacre behind clever description. And there's a real moral difference between blowing up a munitions plant knowing some neighbours will die, and bombing the neighbourhood *in order to* terrorise — the first can be justified, the second never.
- **Example:** Terror-bombing a city to break enemy morale *intends* civilian deaths as the means — DDE forbids it. Striking a legitimate military depot, foreseeing some collateral deaths, can pass — if proportionate.
- **Likely follow-ups:**
 - *Who defends DDE in this unit?* Anscombe and Nagel.
 - *What's the key distinction it rests on?* Intention vs foresight.
 - *What stops it justifying anything?* The proportionality requirement on foreseen harm.
- **If you blank:** DDE permits *foreseen* civilian deaths as a proportionate side-effect but never *intended* ones — so it limits killing, not licenses it.

Should we accept Walzer's "moral equality of combatants"? [Q2]

- **Position:** No — McMahan and Rodin are right to reject it: a soldier fighting an unjust cause is more like an armed robber than a moral equal of his victim.
- **Reason:** Walzer holds that all soldiers have an *equal* right to kill each other regardless of which side's cause is just — once you're in uniform, you and the enemy are moral equals, each permitted to fight by the rules. The revisionists (those rethinking the moral basics) deny this. If your war has no just cause, you are imposing an *unjust threat*; the soldier defending against you is innocent, and you have no more right to kill him than a robber has to shoot the homeowner who resists. Following the in-bello rules doesn't manufacture a right to kill that the injustice of your cause has removed.
- **Objection:** Walzer's view has huge practical advantages: ordinary soldiers can't reliably judge whether their war is just, and granting equal status keeps the laws of war workable and protects POWs — abandon it and every captured enemy is a potential murderer.
- **Reply:** That's a strong reason to keep moral equality as a *legal and practical convention*, but not as the deep *moral* truth. We can say the law should treat soldiers equally for good pragmatic reasons while still recognising that, morally, the unjust combatant does wrong — which matters for how we judge wars and for emerging cases like drones and autonomous weapons.
- **Example:** A soldier conscripted into a war of pure aggression kills a defender repelling the invasion. Walzer says both act permissibly; McMahan says the aggressor's soldier is wrongfully killing an innocent.
- **Likely follow-ups:**
 - *Whose view is "moral equality"?* Walzer's.
 - *Who rejects it and why?* McMahan and Rodin — the unjust combatant is like an armed robber.
 - *Can you keep it legally but reject it morally?* Yes — a workable convention needn't be the deep moral truth.

- **If you blank:** Walzer says all soldiers may kill each other equally; I side with McMahan — the unjust combatant is like an armed robber, not a moral equal.

Define terrorism. Is it always morally worse than conventional war? [Q2]

- **Position:** Terrorism is the deliberate killing of innocents to coerce a wider group, and it's gravely wrong — but Held is right that it isn't *automatically* worse than every conventional war.
- **Reason:** Primoratz defines terrorism as deliberate violence against *innocents* (non-combatants) aimed at terrorising and coercing some wider audience into political concessions. What makes it distinctively wrong is exactly what makes terror-bombing wrong: it *intends* the deaths of the innocent, violating non-combatant immunity at its core. But Held argues we should judge terrorism by the *same* standard as war, not a special harsher one — and once we do, a terrorist act that kills a few may be no worse than a "legitimate" war that foreseeably kills thousands of civilians.
- **Objection:** Surely terrorism is categorically worse: it *targets* the innocent on purpose, whereas conventional war only kills them as a side-effect — that intention-vs-foresight difference makes terrorism a different moral species.
- **Reply:** The intentional targeting is indeed terrorism's signature wrong, and Sparrow and others would hold a robust line there. But Held's point survives: scale and consequences matter too, and a state that knowingly devastates civilian areas can't claim clean hands merely because it didn't "intend" each death. Crucially, that also means *states* can be terrorists — terrorism isn't defined by who does it.
- **Example:** Allied area-bombing of cities in WWII deliberately targeted civilian populations to break morale — on Primoratz's definition that is terrorism, committed by a state, however we usually label it.
- **Likely follow-ups:**
 - *Whose definition?* Primoratz — deliberate violence on innocents to coerce.
 - *Who says judge it like war?* Held — same standard, not automatically worse.
 - *Can states commit terrorism?* Yes — it's defined by the act, not the actor.
- **If you blank:** Terrorism is deliberately killing innocents to coerce (Primoratz); it's gravely wrong, but Held shows it isn't automatically worse than war that kills more.

Can a weapon or act be "mala in se" — evil in itself — and what follows if so? [Q2]

- **Position:** Yes, some means are evil in themselves, and only a deontologist can really say so — which is why the category matters.
- **Reason:** *Mala in se* means acts or weapons wrong *intrinsically* (in themselves), not merely because of their effects — torture, rape, poison, and arguably weapons that cannot discriminate at all. It is a purely deontological verdict (a judgment that the act is wrong whatever it achieves). Here Nagel's clash in "War and Massacre" is decisive: the *utilitarian* judges everything by outcomes, so in principle *any* means is permissible if the results are good enough; the *absolutist* (deontologist) holds that some things you must never do whatever the payoff. Only the absolutist can call something mala in se, because for the pure consequentialist nothing is off-limits in itself — it all depends on the sums.

- **Objection:** Isn't "evil in itself" just disguised consequentialism? We call poison or torture wrong *because* of the horrific effects they tend to cause — so it reduces to outcomes after all.
- **Reply:** No — the test is whether you'd permit it if, hypothetically, the outcome were good. The absolutist says you *still* may not torture an innocent even to avert disaster; that "even if it worked" refusal is precisely what makes the act wrong *in itself* rather than merely usually-harmful.
- **Example:** Designing a weapon to inflict *deliberately hideous* wounds purely to terrify the enemy: even if it "worked" militarily, the absolutist says it's wrong in itself — and that's why such weapons are banned regardless of efficacy.
- **Likely follow-ups:**
 - *Who frames the absolutism/utilitarianism clash?* Nagel, "War and Massacre."
 - *Why can't a consequentialist call anything mala in se?* For them everything depends on outcomes, so nothing is forbidden in itself.
 - *Is this ad bellum or in bello?* In bello — it constrains the *means* of fighting.
- **If you blank:** Some means are evil in themselves (torture, poison); only a deontologist can say so, since for a consequentialist everything depends on outcomes.

Is it permissible to kill someone to stop them taking your chocolate? Where is the threshold for lethal defence? [Q2]

- **Position:** No — you may not kill over chocolate; lethal force is only proportionate against a threat to life or something close to it, and that threshold is the whole point of the exercise.
- **Reason:** This is the Week 2 "escalation ladder": may you kill an armed person who'll take your chocolate? your coat? lock you out of your home? kill you? kill your neighbour? The ladder probes proportionality at the individual level — the same logic that governs war. Killing is the gravest harm you can inflict, so it can only be justified to avert a comparably grave harm. Chocolate, or even a coat, doesn't come close; so even against an *armed* thief, killing would be disproportionate. The line sits where the threat becomes one to life or grievous, irreversible harm.
- **Objection:** But the thief is *armed* and *in the wrong* — he's the aggressor, so by forcing the choice he forfeits his right and you may use whatever force stops him, trivial stakes or not.
- **Reply:** Wrongdoing reduces his protection but doesn't switch it off without limit. Even against an aggressor, your response must be proportionate to what's threatened; otherwise "he started it" would license killing over a parking space. The aggressor's fault widens what you may do, but the *value at stake* still caps it.
- **Example:** Scale it up to war: this is why you can't level a city to stop the theft of a border village — disproportionate at the level of states for exactly the reason it's disproportionate over chocolate.
- **Likely follow-ups:**
 - *Which principle does the ladder test?* Proportionality — at the individual level, mirroring war.
 - *Does the aggressor's fault change things?* It lowers his protection but doesn't remove the proportionality cap.
 - *Where's the threshold?* Roughly: a threat to life or grave, irreversible harm.
- **If you blank:** Killing is only proportionate against a threat to life or grave harm — so no, not over chocolate, however armed the thief.

Q3 BANK — Synthesis / Integration

Apply the principle of distinction (Week 6) to autonomous weapons (Week 11): can a "killer robot" ever satisfy civilian immunity? [Q3]

- **Position:** Even a robot that could mechanically obey distinction should not be allowed to kill, because distinction is not just a sorting rule — it presupposes a responsible agent doing the sorting, and that is exactly what an autonomous weapon removes.
- **Reason:** Distinction (only deliberately target combatants, never the innocent) is the heart of fighting justly, and the best ground for it is responsibility — you may kill someone only if they are responsible for an unjust threat (McMahan). A machine can recognise shapes — a rifle, a uniform — but liability is a moral judgement about a person, not a pattern an object-recogniser can read. And even granting the robot sorts perfectly, Sparrow's responsibility gap bites: when it wrongly kills an innocent, no one — programmer, commander, or machine — can be justly held to account. Distinction without an accountable sorter is hollow.
- **Objection:** Arkin says a machine with an "ethical governor", free of fear and rage, would honour distinction more reliably than a panicked soldier — so it would target fewer innocents and refusing it is what kills more civilians.
- **Reply:** Fewer wrongful deaths is not zero, and each remaining one still has no one to answer for it; consistency on distinction is a consequentialist win, but distinction is a deontological rule about respecting persons, so a tidier killing no human ever weighed is still a wrong. And distinction never travels alone: every strike also needs a proportionality judgement — weighing military gain against collateral civilian harm — which is a moral weighing, not a detection task, and exactly what code can't do.
- **Example:** The August 2021 Kabul strike killed 10 civilians; today we can name and try the humans who approved it. Had an autonomous weapon selected that target alone, the same 10 die but the accountability evaporates — distinction violated, with no one home to answer.
- **Likely follow-ups:**
 - *"If the robot really could judge liability as well as an officer, would you allow it?"* — No, the respect argument still bites: being killed by something that can't grasp your death matters is its own wrong.
 - *"Does a ship's auto-gun shooting incoming missiles violate distinction?"* — No, it targets objects not people, so neither the responsibility nor respect worry applies.
- **If you blank:** Distinction needs a responsible agent to do the targeting; a killer robot removes that agent, so even perfect sorting leaves a wrongful killing no one can answer for.

Is nuclear deterrence (Week 9) defensible on the same grounds the realist (Week 3) defends war? [Q3]

- **Position:** No — deterrence rests on the realist's own argument, but that argument is the prescriptive realism that refutes itself, so deterrence inherits the weakness rather than the strength.

- **Reason:** The realist case for deterrence is consequentialist: in an anarchic world with no world police, a state must threaten mass retaliation to survive, and 80 years without great-power war is the payoff. But this is prescriptive realism — "the state OUGHT to do whatever secures survival" — and the word "ought" is already moral (Orend): the realist has rejoined the moral argument while claiming to escape it. So deterrence can't hide behind "morality doesn't apply." Judged morally, it requires sincerely intending to murder millions of innocents, which violates distinction by design — and only a deontologist can call that wrong-in-itself.
- **Objection:** Descriptive realism is right that survival is non-negotiable, and "ought implies can" — you can't demand a state forgo the one tool that prevents its annihilation, so deterrence is simply forced on it.
- **Reply:** "Ought implies can" limits what morality demands; it doesn't abolish it, and states usually have survivable alternatives short of threatening to incinerate cities. Worse, the consequentialist payoff is shaky: Ward Wilson argues deterrence may not reliably work, and Schlosser documents how near we've come to launch-by-accident. Strip out the unreliable payoff and you're left with the bare wrong of intending mass murder.
- **Example:** The 1983 Petrov incident: a false alarm reported incoming US missiles and doctrine was to retaliate. One officer's refusal is the only reason the "conditional intention" wasn't executed against tens of millions — showing the threat was sincere and the accident risk real.
- **Likely follow-ups:**
 - *"Is there a soft realism that just warns against moralism?"* — Yes, and the just-war theorist should absorb that prudence (proportionality, prospect of success) without accepting the licence.
 - *"If deterrence demonstrably prevented war, would the intention become permissible?"* — On absolutism, no: deliberately intending to murder innocents is off the table whatever the payoff.
- **If you blank:** Deterrence borrows the realist's "survival licenses anything", but that's prescriptive realism, which is self-refuting because "ought" is a moral word — so deterrence is a bad ethics of war, not the absence of one.

Does the double-effect reasoning (Week 8) that excuses some civilian deaths also excuse drone collateral damage (Week 10)? [Q3]

- **Position:** It can excuse genuine drone collateral damage, but only when the strike passes the same tests double effect always demands — and signature strikes fail those tests, so double effect does not launder them.
- **Reason:** The Doctrine of Double Effect (you may cause foreseen-but-unintended civilian deaths as a side-effect of a legitimate strike, never as your means or end) draws its force from the intention/foresight line — a deontological idea. A drone strike on a confirmed, liable target where bystanders die as a regretted side-effect fits the pattern: the deaths are foreseen, not aimed at, and if proportionate and necessary they're excused. But double effect has conditions — the target must be legitimate, the deaths proportionate, the act a last resort — and those are exactly where real drone programmes break.

- **Objection:** A dead child is dead whether "aimed at" or "foreseen", so the intention/foresight line is a convenient fiction that lets powerful states kill civilians with a clear conscience while condemning others.
- **Reply:** The line tracks a real difference: the terrorist needs the innocent dead (their deaths are the plan), the lawful bomber needs the target destroyed (the deaths are a regretted by-product whose absence wouldn't ruin the mission). But the objection is right as a warning — when foreseen deaths become massive and near-certain relative to an uncertain gain, "side-effect" becomes a fig leaf, and that's precisely when proportionality should block the strike.
- **Example:** A confirmed bomb-maker killed with one nearby casualty can fit double effect. The August 2021 Kabul strike — 10 civilians including 7 children, no actual bomb — fails it: the target wasn't liable and the deaths weren't proportionate. Double effect excuses the first, not the second.
- **Likely follow-ups:**
 - *"Why can't a signature strike use double effect?"* — Because it kills people by pattern, not confirmed identity, so the innocent become your target, not a side-effect — the first DDE condition fails.
 - *"At what ratio does 'side-effect' stop excusing?"* — When civilian deaths are large and certain against a speculative gain, proportionality fails even if intention is clean.
- **If you blank:** Double effect can excuse drone collateral deaths only if the target is legitimate, the harm proportionate, and it's a last resort — so it covers a clean strike on a liable target but never a signature strike on a pattern.

Does McMahan's liability account (Week 6) change who a drone may justly target (Week 10)? [Q3]

- **Position:** Yes — liability, not the battlefield or the uniform, is what makes a drone target legitimate, and applying it strictly rules out most signature strikes while permitting a precise strike on a confirmed, responsible individual.
- **Reason:** McMahan's point is that what strips your right not to be killed isn't your clothing or your location but your moral responsibility for an ongoing unjust threat. Carry that into drone war and the licence to kill a named individual depends entirely on whether he is genuinely liable — actively responsible for a serious threat that can't be stopped by arrest. If he is, a precise drone strike that spares bystanders is morally better than a messy raid (Byman). If he isn't, the drone is executing a stranger.
- **Objection:** In real war you fire on enemy positions without checking IDs, so demanding confirmed liability for every strike makes fighting impossible — behaviour is evidence, and a drone reading a "pattern" is just normal battlefield judgement through a camera.
- **Reply:** Context changes what the evidence is worth. In open battle the population has largely separated from fighters, so the conduct inference is reasonable. Over a populated area not in declared war, the same "signatures" — men gathering, carrying tools, attending a funeral — are ordinary civilian life, so the base rate of innocence is high and the inference is weak. Acting on weak evidence to kill is a failure of the precaution liability demands; the camera doesn't fix the inference, it breeds overconfidence in it.

- **Example:** A confirmed al-Qaeda operational planner directing attacks from territory no government controls, uncapturable — liable, so killable. Versus a signature strike on a "military-age male" fitting a movement profile — no confirmed identity, no proven threat — which has hit weddings and rescuers. Liability permits the first, forbids the second.
- **Likely follow-ups:**
 - *"How is a signature strike different from police shooting anyone who 'looks like' a suspect?"* — It isn't, morally — both kill on a profile rather than proven liability, which is why both are wrong.
 - *"Could a civilian ever be liable to a drone strike?"* — In principle yes, if directly responsible for an ongoing threat, but mere support is far below the threshold.
- **If you blank:** McMahan says responsibility for an unjust threat, not the uniform, makes you a target — so a drone may justly kill a confirmed liable individual who can't be captured, but never a stranger picked out by a behaviour pattern.

Does Rodin's self-defence challenge (Week 4) bite harder in riskless drone or robot war (Weeks 10-11)? [Q3]

- **Position:** Yes — Rodin's worry that war-killing borrows the moral force of individual self-defence gets sharper in riskless war, because the self-defence analogy assumes mutual jeopardy that a drone operator or autonomous weapon simply doesn't face.
- **Reason:** Individual self-defence justifies lethal force because you face a lethal threat — your life is genuinely on the line. Rodin already doubts whether national self-defence cleanly inherits that licence. Now strip out the risk entirely: a drone pilot thousands of miles away, or a machine in the field, faces zero danger. If the killing's licence comes from self-defence, and self-defence is grounded in being under threat, then one-sided killing from total safety is the case where the analogy is thinnest — there's no "defence" of the killer at all, just hunting. This is exactly Sparrow's unease with riskless, asymmetric war.
- **Objection:** Force-protection is a virtue, not a vice — we don't tell snipers to step into the open, and a just combatant has no duty to share his enemy's risk; demanding "fair" mutual danger would just get more of your own people killed for no moral gain.
- **Reply:** Granted that minimising your own risk isn't itself wrong — the worry isn't fairness, it's that the self-defence justification quietly disappears when there's no threat to defend against. If the licence to kill doesn't come from self-defence, it must come from the target's liability instead — which is fine for a confirmed liable target but offers no cover at all for killing on suspicion, where most riskless drone and robot war actually operates.
- **Example:** A sniper duel involves mutual jeopardy; a drone operator killing a "signature" target over Waziristan from an air-conditioned trailer involves none. The first looks like combat; the second looks like one-sided killing that has to be justified by the target's liability, not by any defence of the operator — which is why an unconfirmed target is so hard to defend.

- **Likely follow-ups:**

- *"Doesn't this prove all asymmetric war is wrong?"* — No, it shifts the whole justificatory weight onto the target's liability, which a confirmed liable target can still bear.
- *"Does it apply to autonomous weapons too?"* — Even more so: no human faces risk and no human even decides, compounding the self-defence gap with the responsibility gap.
- **If you blank:** Self-defence justifies killing because you're under threat; Rodin doubts states inherit that, and riskless drone/robot war removes the threat to the killer entirely — so the licence must come from the target's liability, not self-defence.

Is the deterrence paradox (Week 9) just the intention/foresight distinction from double effect (Week 8) wearing a different hat? [Q3]

- **Position:** No — they're cousins but not the same: the deterrence paradox is actually harder, because it concerns a sincerely INTENDED wrong, whereas double effect's whole point is to handle merely FORESEEN harms.
- **Reason:** Double effect (Week 8) turns on intention versus foresight: you may foresee civilian deaths as a side-effect but must never intend them as end or means. Deterrence runs into trouble at exactly the intended end of that scale. For MAD to work the threat must be sincere — the enemy must believe you really would annihilate their cities — so deterrence requires a standing, genuine conditional INTENTION to murder millions of innocents. Double effect can't help here, because what's wrong with deterrence isn't a foreseen side-effect; it's the intention itself, which double effect explicitly forbids. So far from dissolving the paradox, double effect sharpens it.
- **Objection:** But the bomb is never meant to be used — you intend it precisely so you never have to act, so this is a conditional intention aimed at no actual deaths, more like a regretted foreseen risk than a true intention to kill.
- **Reply:** A sincere conditional intention is still a real intention, not a bluff — and Kavka's puzzle is exactly that it can be rational to form an intention to do what it would be wrong to do. On a deontological view what you intend is central to what you are doing, so cultivating a genuine intention to slaughter cities is already corrupt, used or not. Double effect's intention/foresight line doesn't rescue you; it convicts you, because the intention is squarely on the forbidden side.
- **Example:** The Petrov false alarm shows the intention was real, not theatrical: doctrine was to retaliate, and only one officer's refusal stopped tens of millions of intended deaths from being carried out. That's an intended-wrong problem, not a foreseen-side-effect problem.
- **Likely follow-ups:**
 - *"So where do they overlap?"* — Both rest on the deontological premise that intention matters morally, not just outcomes — that's the shared engine.
 - *"Could a bluff escape the paradox?"* — Maybe in theory, but a bluff that the enemy might call doesn't deter, so real deterrence needs the real intention.
- **If you blank:** Double effect is about foreseen side-effects; deterrence requires a sincere INTENTION to murder cities — so it's not the same distinction, it's the harder, forbidden side of it, which is why double effect sharpens rather than solves the paradox.

Can Walzer's moral equality of combatants (Week 6) survive autonomous weapons (Week 11)? [Q3]

- **Position:** No — autonomous weapons make Walzer's moral equality of combatants almost unintelligible, because there is no combatant on the machine's side to be morally equal to anyone, which pushes the whole debate toward the revisionist, liability-based picture.
- **Reason:** Walzer's moral equality of combatants says every soldier may kill enemy soldiers regardless of whose cause is just, partly because each shares the risk and mutual jeopardy of battle. McMahan already rejects this — an unjust combatant is like an armed robber, not a moral equal. Autonomous weapons strip the idea bare: a machine isn't a combatant with rights and risks, it's a tool with no moral standing, no jeopardy, nothing to be "equal" about. The human who deployed it isn't on the battlefield either. So there's no symmetric warrior relationship left — only a question of whether the human victim is liable, which is precisely McMahan's frame.
- **Objection:** Walzer's principle was never about literal symmetry of weapons — armies have always been unequal (tanks vs rifles) — it's a convention that lets soldiers fight without each being a murderer, and that convention can survive whatever hardware is involved.
- **Reply:** The convention rests on combatants being agents who accept risk and can be held responsible; a machine breaks both. You can't extend "you may kill me and I may kill you, and neither is murder" to something that can't be wronged, can't be held responsible, and faces no risk. So the convention doesn't just get more lopsided — it loses the agents it was a convention between, which is why the case naturally collapses into asking about liability and Sparrow's responsibility gap.
- **Example:** Picture an autonomous turret versus human soldiers. We don't ask "is the turret morally equal to the soldiers?" — the question is nonsense. We ask whether those it kills are liable and who answers for its mistakes. That's the revisionist/liability question, not Walzer's equality.
- **Likely follow-ups:**
 - *"Does this vindicate McMahan?"* — Largely yes: once the symmetric-warrior picture breaks, liability is the only thing left doing moral work.
 - *"Could two autonomous armies fight each other 'equally'?"* — Then no persons are at risk at all, which raises whether it's even war or just machines breaking machines.
- **If you blank:** Moral equality of combatants assumes two risk-bearing human soldiers; an autonomous weapon is a tool with no standing and no jeopardy, so there's no equality to speak of — the question collapses into McMahan's liability and Sparrow's responsibility gap.

Is terrorism (Week 7) morally worse than a signature drone strike (Week 10)? [Q3]

- **Position:** Terrorism is worse per act because it deliberately targets the innocent, but a signature strike is closer to terrorism than states like to admit — and on a consistent, act-based standard some signature strikes shade into the same wrong.

- **Reason:** Primoratz defines terrorism as deliberately using violence against innocents to terrorise and coerce — the innocent are the target, not a side-effect, which is the gravest jus-in-bello wrong. A signature strike aims at a perceived threat, so in intention it's not classic terrorism. BUT a signature strike knowingly kills people whose identity and liability are unconfirmed — it accepts a real chance of killing innocents AS the target, by pattern, not as a foreseen side-effect of hitting a confirmed combatant. That's distinction breaking at the root, which is uncomfortably close to terrorism's core wrong.
- **Objection:** This collapses a vital line — the drone operator is trying to hit a threat and regrets civilian deaths, while the terrorist wants civilians dead; intention separates them, so calling signature strikes "near-terrorism" is a false equivalence (and Held's "judge both the same" overreaches).
- **Reply:** Held is right that we shouldn't run a double standard that condemns the weak's violence while excusing the state's. The intention/foresight line does real work for a confirmed target — but a signature strike's "target" is itself just a probability, so the operator isn't cleanly aiming at a known combatant; he's accepting innocent deaths as the price of killing a pattern. When the pattern is ordinary civilian life, that recklessness about who you kill starts to resemble deliberately treating the innocent as expendable.
- **Example:** 9/11 chose civilians precisely because killing them spreads fear — paradigm terrorism. A "double-tap" signature strike that hits rescuers returning to a blast site kills people by pattern with reckless disregard for their innocence — not identical to 9/11, but on an act-based, no-double-standard view it's far closer to it than to a clean strike on a confirmed planner.
- **Likely follow-ups:**
 - *"Can states commit terrorism?"* — Yes — on an act-based definition (Coady), deliberately terrorising civilians is terrorism whoever does it, uniform or not.
 - *"So is a confirmed-target drone strike terrorism?"* — No, that's a legitimate strike under distinction and double effect; the problem is the signature, not the drone.
- **If you blank:** Terrorism deliberately targets the innocent (Primoratz), so it's worse per act than a strike aimed at a threat — but a signature strike kills by pattern, not confirmed identity, so on a consistent act-based standard it edges toward the same wrong of treating the innocent as expendable.

Does Sparrow's "robots and respect" argument (Week 11) also condemn riskless drone killing (Week 10)? [Q3]

- **Position:** Partly — the respect argument bites hardest against autonomous weapons, but a weaker version reaches riskless drone killing too, though a human at the trigger keeps it on the right side of the line.

- **Reason:** Sparrow's "robots and respect" says being killed by a machine that decides your death fails to respect your moral status, because no human agent ever weighs the gravity of ending your life — you're processed, not judged. The thing that does the moral work is whether a human registers that your death matters. A drone keeps a human at the trigger, so someone does decide and could be held to account — the respect condition is met, even from total safety. That's the crucial difference between a drone (human decides) and an autonomous weapon (machine decides).
- **Objection:** But a drone operator killing a "signature" target from a trailer, watching a screen, treating people as heat-signatures and patterns, barely "respects" the victim more than a robot would — so if respect condemns the robot, it should condemn riskless drone killing too.
- **Reply:** There's a real residue here — remote, riskless killing-by-pattern does erode the respect a genuine human judgement embodies, which is why Sparrow is also a sceptic of riskless asymmetric war. But the drone still has a human who can weigh the decision and answer for it, so the respect failure is partial, not total; with an autonomous weapon it's complete — no one weighs anything. So the argument condemns signature-strike drone killing as a degraded case while reserving its full force for the machine that removes the human entirely.
- **Example:** The Kabul strike: humans on screens, far away, misread a man as a threat and killed 10 — respect was thin and the judgement was bad, but a human did decide and can be held to account. An autonomous weapon making the same kill would remove even that thin thread of human judgement — the difference between a degraded respect and none at all.
- **Likely follow-ups:**
 - *"Doesn't respect also condemn high-altitude bombing, where the killer never sees the victim?"* — It strains those cases too, which is why a human deciding to weigh each strike matters more than physical proximity.
 - *"If a human just rubber-stamps the machine's choice, is respect satisfied?"* — No — nominal human-on-the-loop without genuine judgement is the machine deciding in disguise.
- **If you blank:** Sparrow says being killed by a machine that decides your death disrespects you because no human weighs it; a drone keeps a human at the trigger so the respect condition is (barely) met, but riskless killing-by-pattern erodes it — full force lands on the autonomous weapon, partial force on the signature strike.

How does the consequentialism-vs-deontology clash (the Master Key) explain BOTH the Arkin-vs-Sparrow robot debate (Week 11) AND the nuclear deterrence debate (Week 9)? [Q3]

- **Position:** Both debates are the same Level-0 clash in different costumes — consequentialism (only outcomes matter) versus deontology (some acts are wrong in themselves) — and in both I side with deontology, because killing is exactly the case where outcomes can't be the whole story.

- **Reason:** Arkin's case for killer robots is pure consequentialism: if a machine with an ethical governor kills fewer civilians than a frightened soldier (think My Lai), then outcomes favour the robot and banning it is what kills more. Sparrow answers deontologically: fewer wrongful deaths still leaves each one with no one accountable (responsibility gap) and no human respecting that it mattered (respect) — wrongs in the act, not the tally. Nuclear deterrence is the identical structure: the defender says the good outcome (80 years of peace) licenses the threat; the critic says sincerely intending to murder millions is wrong in itself whatever the payoff. Same engine: net outcomes versus inviolable side-constraints.
- **Objection:** If you keep choosing the deontological side, you're choosing worse outcomes — more dead civilians from soldiers over robots, and risking nuclear war by refusing deterrence — which looks like keeping your hands clean at others' expense.
- **Reply:** Two replies that work in both cases. First, the consequentialist premises are shaky: machines can't make proportionality judgements (a moral weighing, not a detection task), and Ward Wilson and Schlosser show deterrence may not work and risks catastrophe-by-accident — so the promised good outcomes may be illusory. Second, even granting the outcomes, some methods of killing are off-limits the way torture is — we already accept that good consequences don't license every means. Killing where no one answers, or intending mass murder, are two such off-limits methods.
- **Example:** My Lai for the robot debate (a human atrocity that was named, tried, reckoned with — vs an autonomous massacre no one answers for) and the Petrov false alarm for deterrence (a sincere intention to kill tens of millions, stopped only by one man's refusal). Each shows the deontological worry — accountability, intention — surviving even a favourable body count.
- **Likely follow-ups:**
 - *"So is the whole unit just this one distinction?"* — Largely yes — almost every debate reduces to outcomes-only vs acts-wrong-in-themselves, with tools like intention/foresight and rights hanging off it.
 - *"When does consequentialism win?"* — Inside the side-constraints — proportionality and prospect-of-success are consequentialist tests JWT keeps, but capped by absolute bans like never targeting the innocent.
- **If you blank:** Arkin-vs-Sparrow and the deterrence debate are both consequentialism (outcomes only) vs deontology (some acts wrong in themselves) — robots/cities may give better numbers, but a killing no one answers for and an intention to murder millions are wrongs in the act, so I side with deontology in both.