

Deliverable B — The Confident Bluffer's Toolkit

This is the reusable machinery. None of it requires knowing the answer cold. It works by leaning on the rubric: a clear structure, a stated opinion, and one reason will out-score a rambling expert every time. Drill these until they're automatic, then you can walk into any of the three questions and *sound* fluent.

1. The Universal Answer Template

Almost every question in this unit is, underneath, the same fight: **do only outcomes matter (consequentialism), or are some acts wrong in themselves whatever the outcome (deontology)?** That means you can answer *almost anything* by running this one move:

"There are broadly two ways to look at this — a consequentialist line and a deontological line. The consequentialist would say [whatever produces the best overall outcome — fewest deaths, most lives saved — is what's permitted]. The deontologist would object that [some acts are forbidden in themselves — you may never *deliberately* target the innocent, even to win]. **My own view is** [pick one — say it in a single sentence] **because** [give your one reason, then name the strongest objection and bat it away with one line]."

This bracketed skeleton is the **core** of an answer, not the whole thing. In the room you stretch it toward the house target — roughly 350–500 words, 3–4 minutes — by *lingering*: spell the reason out, give the objection its strongest form, and finish with one concrete example. The skeleton guarantees you never lose the thread; the lingering fills the time.

Filled in — on autonomous weapons ("should we build killer robots if they'd kill fewer civilians?")

"There are broadly two ways to look at this. The consequentialist line — Arkin's — says that if a robot with an 'ethical governor' feels no fear or rage and so kills *fewer* civilians than a stressed soldier, then building it is the humane choice and *banning* it is what's immoral. The deontologist — Sparrow's — objects that good outcomes don't settle it: even if the robot kills fewer people, there's still no one who can be justly held to account for each wrongful death, and being killed by a machine that can't grasp that your death *matters* fails to respect you as a person. My own view is that we should not build them, because killing is exactly the case where accountability and respect aren't optional extras — a tidier body count doesn't fill the responsibility gap. The strongest objection to my view is that I'm choosing more dead civilians for the sake of a philosophical scruple; but I'd reply that we already accept some methods are off-limits regardless of outcome — that's why we ban torture even when it might extract life-saving information."

That's the **compressible core** — about two minutes spoken. To hit the full 3–4 minutes, you expand it: name Sparrow's two distinct arguments explicitly (the *responsibility gap* and "*Robots and Respect*"), add the My Lai example to make Arkin's case vivid, and close on the framing line — "it's consequentialism versus deontology in a robot costume, and I come down on the deontology side." You can swap the bracketed content for nuclear weapons, terrorism, targeted killing, mala in se — the *skeleton* never changes.

Why this auto-satisfies the rubric

- **Criterion 2 (clarity of explanation):** the "two ways to look at this" frame *is* a structure. The examiner can hear your scaffolding, so the explanation lands as organised, not rambling.
- **Criterion 3 (taking a position):** the phrase "**My own view is...**" forces you to land a stance every single time. You physically cannot finish the template without an opinion.
- **Criterion 4 (supporting with argument):** "**because...**" forces a reason, and "**the strongest objection... but I'd reply...**" delivers reasoned *back-and-forth*, which is exactly what "supporting that opinion with reasoned argument" means. You've now argued *both* directions and still chosen.

Three of the four criteria (the ~75% that's about form) are satisfied by the template alone, before you've said anything clever about the actual topic.

2. Signposting Phrases (buy thinking time, sound scholarly)

Say one of these *out loud* while your brain catches up. They are not filler — each one frames the question and earns clarity marks.

Stall (start talking before you know the ending)

1. "That's a good question — and I think the first thing to get clear on is what exactly we're asking..."
2. "There's a lot packed into that, so let me pull it apart."
3. "It's worth slowing down on the key term here, because a lot turns on how we define it."
4. "Let me start with the strongest version of the opposing view, then say where I come down."

Structure (organise as you go — pure clarity marks)

5. "It's worth distinguishing *jus ad bellum* — whether it was just to go to war — from *jus in bello* — whether it's being fought justly. This question is really about the second."
6. "There are broadly two camps on this..."
7. "Let me separate two things people often run together here..."
8. "So there's a question of principle and a separate question of practice, and I'd answer them differently."

Take a stance (this is where Criterion 3 lives — never skip it)

9. "Where I come down is this:..."
10. "My own view is that..., and here's the one reason that does the work for me..."

- .1. "If I had to plant a flag, I'd say..."
- .2. "That really turns on whether you think *outcomes* or *duties* matter more — and I think duties win in this case, because..."

Handle a follow-up (when they push back)

- .3. "The strongest objection to my own view would be..."
 - .4. "I'll happily grant that point — but I don't think it defeats my position, because..."
 - .5. "That's exactly the hard case, and I think it sharpens my view rather than sinking it."
 - .6. "Fair — let me bite the bullet there: I'd accept that consequence and here's why I can live with it..."
-

3. Name-Drop Bank

Attribute these *exactly*. Getting a name right scores Criterion 1 (knowledge); swapping two names is the one error that looks worst. One name dropped correctly per answer is plenty.

Thinker	Known for (one line)	The exact claim to attach
Walzer	The modern anchor of Just War Theory	Self-defence against aggression ("legalist paradigm"); the moral equality of combatants — every soldier may kill enemy soldiers regardless of whose cause is just
McMahan	Revisionist; the liability view	What makes you a legitimate target is moral responsibility for an unjust threat , not your uniform — so he <i>rejects</i> moral equality of combatants (an unjust soldier is like an armed robber)
Rodin	Reductivist challenge to self-defence	National self-defence is not cleanly analogous to individual self-defence — individual defence protects a <i>life</i> , but resisting a bloodless conquest only protects <i>institutions</i>
Primoratz	Defines terrorism; immunity theorist	Terrorism = deliberately using violence against innocents to terrorise and coerce a wider audience; grounds civilian immunity in responsibility
Held	Terrorism vs war	Judge terrorism by the same standard as conventional war — it's not <i>automatically</i> worse; condemning the weak's violence while excusing the state's is a double standard

Thinker	Known for (one line)	The exact claim to attach
Sparrow (<i>your examiner</i>)	Autonomous weapons; robust civilian immunity	Two arguments against killer robots: the responsibility gap (no one — programmer, commander, or machine — can be justly held to account) and "Robots and Respect" (being killed by a machine fails to respect your moral status). Also W6 robust immunity ("Hands up who wants to die?") and sceptic of riskless/asymmetric war
Nagel	"War and Massacre"	Absolutism vs utilitarianism : some acts (deliberately attacking the innocent) are forbidden whatever the payoff; limits on <i>whom</i> and <i>how</i> you may attack
Anscombe	"War and Murder"	Anti-pacifist, pro just-war + Double Effect ; pacifism is <i>pernicious</i> because branding all killing "murder" erases the combatant/innocent line
Arkin	"The Case for Ethical Autonomy"	PRO autonomous weapons — a machine with an "ethical governor" feels no fear/rage and could obey the laws of war <i>better</i> than soldiers (cf. My Lai); banning them would be the immoral choice
Byman	"Why Drones Work"	PRO drones — most precise tool available; kill <i>fewer</i> civilians than air strikes or ground invasions, and protect your own troops
Boyle	The anti-drone case	ANTI drones — strategically self-defeating; civilian deaths and being hunted from the sky radicalise populations (strategic blowback)
Ward Wilson	"The Gordian Knot"	Doubts deterrence works; the atomic bombs may not have caused Japan's surrender — the Soviet invasion plausibly did
Morgenthau	Founding realist	In an anarchic system states act on national interest as power , not morality — but he still leaves the statesman <i>prudent</i> duties (which is why Orend says realists "smuggle morality back in")
Tolstoy	Absolute pacifism	War is never justified; patriotism is the engine that manufactures the loyalty and hatred mass killing needs

Thinker	Known for (one line)	The exact claim to attach
Annan	R2P	"Sovereignty as responsibility" — sovereignty is <i>conditional</i> on a state protecting its people; if it commits atrocities, its sovereignty no longer shields it (Responsibility to Protect)
Coady	Critic of humanitarian war	Humanitarian intervention is dangerously open to abuse — strong states dress self-interest as rescue and overrate their own good effects; tighten the conditions, don't abandon the principle

4. Pivot Moves

When a question lands on a sub-topic you don't know, *concede the framing for one sentence*, then pivot to ground you own. Done right it sounds like deepening, not dodging.

1. Any drone / targeted-killing question you're shaky on → pivot to asymmetry & liability.

"What's really doing the moral work here isn't the drone technology — it's two questions: is the target genuinely liable in McMahan's sense, and is killing them a last resort? The technology just makes it easier to drop those safeguards. Sparrow's worry about riskless, one-sided killing is the version of this I find most interesting..."

2. Any *jus in bello* / "is this attack OK?" question → pivot to distinction. *"Underneath this is the principle of distinction — you may only ever deliberately target combatants, never non-combatants. So the real question becomes why civilians are immune: is it because they're innocent, harmless, or not responsible for the threat? I'd argue it's non-responsibility — and that's exactly Sparrow's robust-immunity line..."*

3. Any nuclear / weapons question → pivot to mala in se + the master key. *"This is really the mala in se question — is the weapon evil in its very nature, not just in its effects? And that only makes sense if you're a deontologist rather than a consequentialist, which is the fault line running under the whole unit. A strategic nuke can't discriminate and can't be contained, so..."*

4. Any vague "is war ever justified / who's right about war" question → pivot to the three traditions. *"There are only really three stances on offer — Realism (morality doesn't apply to war), Pacifism (war is never justified), and Just War Theory (sometimes, under strict criteria). I think JWT is the defensible middle, because Realism over-reaches one way — we plainly do condemn Bucha — and Pacifism over-reaches the other..."*

5. Any thinker / historical-case question you can't place → pivot to the *ad bellum* / *in bello* split. "The first thing I'd want to do with this case is separate two questions people constantly run together: was it just to go to war, and was it fought justly? They're independent — a war can be right to start and full of atrocities, or wrong to start and fought cleanly. Take Iraq 2003 versus Abu Ghraib..."

5. "If You Remember Nothing Else" — One Line Per Week

If you blank completely, say the line for that week. Each one states a position *and* a reason, so each one scores on its own.

- **Week 1 (Foundations):** "The unit rests on three things — we reject relativism (we really *can* judge My Lai wrong), there are three stances (Realism, Pacifism, Just War Theory, and JWT is the defensible middle), and JWT splits into two independent halves: *jus ad bellum* (just to go to war?) and *jus in bello* (fought justly?)."
- **Week 2 (Pacifism):** "Distinguish *absolute* pacifism (killing is wrong in itself — Tolstoy) from *contingent* pacifism (war never pays off in practice); the contingent version is stronger, and Anscombe's killer objection is that pacifism is *pernicious* because calling all killing 'murder' erases the combatant/innocent line."
- **Week 3 (Realism):** "Split realism in two — *descriptive* realism (states *do* act on interest in anarchy) is often true and a useful warning; *prescriptive* realism (states *should* ignore morality) is self-refuting, because the word 'should' is already moral and Nuremberg shows we really do judge wars."
- **Week 4 (Jus ad bellum):** "The standard justification for going to war is self-defence against aggression (Walzer), but Rodin's challenge is the sharpest objection — national self-defence can't be grounded like individual self-defence, because individual defence protects a *life* while resisting a bloodless conquest only protects *institutions*."
- **Week 5 (Humanitarian war):** "My line is 'sovereignty is no shield for atrocity' (Annan's R2P) — a state's right to be left alone is conditional on protecting its people — but because the doctrine is so easily abused and so often backfires (Coady; Libya), the just cause is real while the bar for actually intervening must be very high."
- **Week 6 (Distinction):** "You may only ever *deliberately* target combatants; the deep question is *why* civilians are immune, and the best answer is that they're *not responsible* for the unjust threat — which is why we hold a robust, near-absolute line (Sparrow) rather than letting immunity erode as civilians' war-support grows."
- **Week 7 (Terrorism):** "Terrorism = deliberately using violence against innocents to terrorise and coerce a wider audience (Primoratz); it's wrong because it makes the innocent the *target*, not a side-effect — and on a consistent, act-based definition, *states* can do it too (Held, Coady)."

- **Week 8 (Mala in se):** "*Jus in bello* isn't just distinction and proportionality — it also bans *mala in se*, means evil in themselves like torture, rape and poison, which are wrong however good the cause; and only a deontologist (Nagel's absolutist), not a consequentialist, can explain why."
- **Week 9 (Nuclear weapons):** "Strategic nukes break every *in bello* rule at once — they can't discriminate and can't be contained — so they're arguably *mala in se*; and they add the deterrence paradox: it may be wrong to *threaten* what it would be wrong to *do*."
- **Week 10 (Targeted killing):** "Targeted killing blurs war (kill combatants, no trial) and law enforcement (arrest, don't execute), so it's only defensible when the named target is genuinely *liable* and capture is impossible — which is exactly why *signature strikes* (killing by pattern, not confirmed identity) are the hardest to justify."
- **Week 11 (Autonomous weapons — Sparrow's own field):** "An autonomous weapon *selects and engages* targets with no human pulling the trigger, and Sparrow gives two reasons to ban it — the **responsibility gap** (no one can be justly held to account for a wrongful killing) and '**Robots and Respect**' (being killed by a machine fails to respect you as a person); Arkin replies that robots could kill *fewer* civilians, but good outcomes don't fill the gap — it's consequentialism vs deontology in a robot costume, and I side with deontology."