

Practice Questions + Talking Points

A big bank of likely questions for each of your three exam questions, with dot-point talking points to riff off.

How to use: pick a question, glance at the points, then talk for ~3–4 minutes out loud. Each set of points runs roughly: **a stance to open with** → **the main reason** → **the other side** → **your reply** → **an example** → **a term/thinker to drop**. You don't need every bullet — take a clear position, give a reason, and you're scoring. There's no "right" answer.

Q1 — The modern hard cases (weeks 9–11) ★

One of your three questions is on these three topics. The biggest focus.

Nuclear weapons

1. Are nuclear weapons mala in se -- evil in themselves -- or just very destructive?

- Stance: I think strategic nuclear weapons are arguably mala in se -- not just powerful, but wrong in themselves to use.
- Main reason: a weapon is mala in se when you literally cannot use it without doing something intrinsically forbidden -- and you can't drop a city-busting nuke without killing huge numbers of innocent civilians, so it fails distinction and proportionality by its very nature.
- Contrast with 'just destructive': a thousand-pound bomb is terrible but in principle you can aim it at a tank column; a strategic nuke's blast, fire and radiation can't be contained to combatants -- the indiscriminateness is baked in.
- Other side: a defender might say it's the USE that's wrong, not the object -- a tiny tactical nuke on an isolated fleet at sea might in theory hit only combatants, so the weapon isn't intrinsically evil.
- Reply: even granting that edge case, the weapons that actually exist and matter strategically are city-killers, and the radiation lingers for generations and crosses borders, so realistically they can never satisfy the in bello rules.
- Term/thinker to drop: this is a deontological verdict -- mala in se sits alongside poison gas and torture, things Nagel in 'War and Massacre' would say absolutism forbids whatever the outcome.
- Closing line: so my view is they're best understood as intrinsically unusable within just war theory - - evil in themselves, not merely big.

2. Can it be ethical to THREATEN to do something it would be unethical to actually do? (the deterrence paradox)

- Stance: this is the deterrence paradox, and my view is the threat is tainted by the act -- if actually incinerating cities is murder, then sincerely intending to do it is also morally compromised.

- Main reason: deterrence only works if the threat is sincere -- you have to genuinely intend to retaliate, and forming a real intention to do something monstrous is itself a moral act, not a free pass.
- The puzzle stated plainly: that's exactly Reflective Exercise 9.2 -- how can it be wrong to DO X but fine to form a firm conditional intention to do X?
- Other side: the consequentialist says judge it by results -- if the threat is never carried out and it keeps the peace for decades, then threatening saved millions of lives, so the intention is justified by what it prevents.
- Reply: but you're either bluffing -- in which case it's not real deterrence -- or you genuinely mean it, in which case you're a leader prepared to murder millions, and that conditional intention corrupts you whether or not the button is ever pushed.
- Term/thinker to drop: this is the conditional-intention problem -- Kavka and Gauthier territory -- and it cuts cleanly along consequentialist (outcomes) versus deontological (the intention itself matters) lines.
- Closing line: so I'd say deterrence forces you to choose between being ineffective and being willing to commit mass murder -- and that's a deeply uncomfortable place to stand.

3. Does nuclear deterrence even work -- and does it matter morally whether it does?

- Stance: I'm sceptical that deterrence works as cleanly as the standard story claims -- and that scepticism really weakens the moral case for keeping nukes.
- Main reason: the whole justification for nuclear weapons is consequentialist -- 'they keep the peace' -- so if they DON'T actually deter, the main argument for tolerating these mala in se weapons collapses.
- Key thinker: Ward Wilson in 'The Gordian Knot' challenges the standard story -- he argues nuclear weapons may not really deter, and that it was the SOVIET invasion, not Hiroshima, that actually forced Japan's surrender.
- Why that bites: if even Hiroshima -- the textbook case -- didn't do the work we credit it with, then the empirical foundation of deterrence theory is shakier than everyone assumes.
- Other side: defenders point to the long peace -- no great-power war since 1945, no nukes used in anger -- and say that track record is no accident, MAD made superpowers cautious.
- Reply plus risk: correlation isn't proof, and Schlosser on 'World War Three by Mistake' shows how many near-accidents there have been -- deterrence also carries a permanent risk of catastrophe by error.
- Closing line: so morally it matters enormously whether it works -- if the consequentialist payoff is illusory, you're left holding an intrinsically wrongful weapon for nothing.

4. Was the bombing of Hiroshima ever justifiable?

- Stance: I lean towards saying Hiroshima was NOT justifiable -- it deliberately targeted a city full of civilians, which is exactly what jus in bello forbids.

- Main reason: distinction and non-combatant immunity say you may only target combatants, and a city is overwhelmingly non-combatants -- so this was the intentional mass killing of innocents, not collateral damage.
- The standard defence: the consequentialist case is that the bomb ended the war fast and prevented a bloody invasion that would have cost far more lives on both sides -- fewer deaths overall.
- Reply on the facts: Ward Wilson disputes even this -- he argues the Soviet entry into the war, not the bomb, actually triggered the surrender, so the lives-saved calculation may rest on a myth.
- Reply on principle: even if it did save lives, an absolutist like Nagel would say you still can't deliberately incinerate a hundred thousand civilians to achieve it -- some acts are off the table whatever the maths.
- Thinker to drop: Glover in 'Humanity' treats Hiroshima as a moral turning point -- the moment restraints on killing the innocent eroded -- which frames it as a fall, not a triumph.
- Closing line: so my honest view is it might have been militarily effective and still morally wrong -- effectiveness doesn't launder the deliberate killing of the innocent.

5. Is it worse to die of radiation poisoning than of blood poisoning from shrapnel? (Reflective Exercise 9.1)

- Stance: my gut says radiation death feels specially horrifying, but when I think it through I'm not sure that intuition tracks a real moral difference.
- The probe's point: this is Reflective Exercise 9.1, and it's testing whether there's something UNIQUELY wrong about nuclear death or whether we're just scope-blind and spooked by the unfamiliar.
- Case for a difference: radiation is invisible, lingering, it harms unborn generations and contaminates land for decades -- it keeps killing long after the war ends, which shrapnel doesn't.
- Case against a difference: dead is dead -- a soldier bleeding out from an infected shrapnel wound suffers just as much, so maybe we're inconsistent, recoiling at the new horror while accepting the familiar one.
- Where it actually bites: the real moral wrong of nukes isn't the MANNER of death but WHO dies -- the indiscriminate killing of non-combatants -- so focusing on radiation-versus-shrapnel can distract from the distinction problem.
- Term to drop: this connects to mala in se -- if radiation's uncontrollable, cross-generational reach is what makes the weapon impossible to use proportionately, then the manner of death does feed back into intrinsic wrongness.
- Closing line: so I'd say the special horror is partly real -- the lingering, borderless harm -- but the deeper wrong is still that nukes can't tell combatant from child.

6. Do tactical nuclear weapons change the moral picture, or are all nukes the same?

- Stance: tactical nukes complicate the easy 'all nukes are mala in se' line, but I still think the moral picture stays grim.

- The distinction: strategic nukes are city-busters meant for deterrence and mass destruction; tactical nukes are smaller, battlefield weapons -- so in principle a tactical nuke aimed at a tank division in a desert might mostly hit combatants.
- Why that matters: if a weapon CAN in principle satisfy distinction and proportionality, it's harder to call it intrinsically evil -- the wrong would be in the use, not the object, so the mala in se verdict weakens for the small end.
- Other side -- the escalation problem: the real danger is that using ANY nuke shatters the nuclear taboo and invites tit-for-tat escalation up to strategic exchange, so even a 'clean' tactical strike risks the unthinkable.
- Reply: that's a partly consequentialist worry about thresholds -- lowering the bar to first use -- and it means even militarily 'usable' nukes carry a catastrophic tail risk that a conventional weapon doesn't.
- Term to drop: Quester is relevant here on whether you could substitute conventional for nuclear weapons -- if you can achieve the same military aim conventionally, the proportionality case for ANY nuke gets much harder to make.
- Closing line: so tactical nukes blur the mala in se line, but the escalation risk drags the moral verdict right back towards 'don't'.

7. Is keeping a nuclear arsenal compatible with just war theory at all?

- Stance: I think it sits in deep tension with just war theory -- you're maintaining weapons you arguably could never justly USE.
- Main reason: just war's in bello rules require distinction and proportionality, and strategic nukes fail both by design, so a JWT purist has to say the arsenal is built around an unusable weapon.
- The deterrence wrinkle: but the whole point isn't to use them -- it's to threaten, to deter -- which is why this collapses straight into the deterrence paradox of intending what you must never do.
- Other side: a realist like Morgenthau would say this is naive -- states exist in anarchy, security is survival, and you keep nukes because rivals have them, morality of war theory doesn't get a vote.
- Reply: that's descriptive realism, not a justification -- Orend would push back that 'everyone does it' doesn't make it right, and JWT is exactly the framework for judging these choices morally.
- Honest middle: many theorists land on a reluctant, conditional toleration -- deterrence as a temporary necessary evil pending disarmament -- which is more an uneasy truce with JWT than a clean fit.
- Closing line: so my view is the arsenal is at best a deeply uncomfortable exception JWT tolerates rather than endorses -- and that discomfort should push towards reduction, not complacency.

8. If launching nukes is wrong, is the leader who builds a deterrent doing something wrong right now -- even in peacetime?

- Stance: I'd argue yes -- if the act would be wrong, a leader who genuinely intends to do it under conditions is already morally implicated, even before anyone fires.

- Main reason: deterrence requires a SINCERE conditional intention -- the leader has to truly be prepared to incinerate millions if attacked -- and forming that settled intention is itself a morally serious act, not a neutral policy.
- Plain version: a person who genuinely resolves 'if you hit me, I will murder your children' is already in a corrupt moral state, even if the trigger condition never arrives.
- Other side: the consequentialist replies that what matters is outcomes, and if the threat keeps the peace and is never executed, then the standing intention has produced only good -- so blame is misplaced.
- Reply: but you can't have the good outcome WITHOUT the genuine willingness -- a known bluff doesn't deter -- so the peace is purchased precisely by leaders being authentically ready to do the monstrous thing.
- Term/thinker to drop: this is the conditional-intention problem again -- Kavka's puzzle -- and it shows deterrence can't separate clean hands from dirty intentions the way defenders want.
- Closing line: so my answer is that nuclear deterrence implicates its wielders morally in peacetime - - the wrong of the act seeps back into the standing threat that makes it work.

Drones & targeted killing

9. Is targeted killing by drone more like an act of war, or more like an extrajudicial execution?

Where do you land?

- My view: it sits in an uncomfortable grey zone, but most drone strikes are best understood as closer to execution than war, because they usually happen OUTSIDE a declared armed conflict against a named individual.
- The core distinction: in war you're allowed to kill enemy combatants on sight; in law enforcement you must arrest and try a criminal, not execute them. Targeted killing collapses that line.
- Why it matters which one it is: if it's 'war', the just war rules (distinction, proportionality) apply and killing combatants is legitimate; if it's 'crime', you owe the target due process and you've skipped the courts.
- The other side: defenders say modern terrorists ARE combatants in a real, ongoing armed conflict, just not one with front lines or uniforms, so it genuinely is warfare and the war rules apply.
- My reply: even granting that, the fact it happens far from any battlefield, against someone who poses no immediate threat at the moment of the strike, makes it look much more like a premeditated killing of a specific person, which is what 'execution' means.
- Definition to nail: targeted killing = the intentional, premeditated killing of a specific named individual by a state, typically outside a formally declared armed conflict, increasingly via drones (Coll's 'The Unblinking Stare' documents this reality in Pakistan).
- Closing line: the deepest problem is that the technology lets states have it both ways, claiming the permissions of war while dodging the accountability of the courtroom.

10. What, if anything, is morally wrong with so-called signature strikes?

- My view: signature strikes are the most ethically indefensible part of the whole drone programme, because you're killing people whose identity you don't even know.

- First, the definition: a signature strike targets people based on a 'pattern of life' or pattern of behaviour, military-age males meeting in a certain way, carrying weapons, rather than a confirmed identity, unlike a 'personality strike' on a named individual.
- The main problem: the principle of distinction requires you to be sure your target is a legitimate combatant, and a behavioural 'signature' is just a guess that can sweep up wedding parties, tribal gatherings, or men carrying rifles for ordinary reasons in that region.
- It also wrecks the combatant/criminal question from the other angle: you can't even ask whether they should be arrested rather than killed if you don't know who they are.
- The other side: defenders argue that in practice you can almost never get a confirmed ID on a battlefield, and that the behavioural patterns are reliable enough to count as evidence someone is part of an armed group.
- My reply: 'reliable enough' is a consequentialist gamble with innocent lives, and the burden of proof in killing should be on certainty, not on probability, which is exactly what distinction demands.
- Thinker to drop: this connects to McMahan on liability, what makes you a legitimate target is your moral responsibility for an unjust threat, not your uniform or your demographic, and a 'signature' tells you almost nothing about responsibility.
- Closing line: if you wouldn't convict someone in a court on a 'pattern of life', you certainly shouldn't kill them on one.

11. Drone warfare is radically asymmetric, the operator is in zero physical danger. Does that riskless quality raise a distinct moral problem?

- My view: yes, there's something genuinely troubling about riskless killing, and I'd side with Sparrow's scepticism here rather than dismissing the worry.
- The plain version of the worry: warfare has traditionally involved mutual risk, two sides both putting themselves in harm's way; when one side faces zero danger, it starts to look less like a fight and more like an execution or even a hunt.
- Why this might be a real moral problem: the willingness to risk your own life is part of what has historically justified taking life in war, and the courage and sacrifice that ennobled the soldier disappears when you're killing from an air-conditioned trailer.
- It also lowers the threshold for war: if no body bags come home, leaders find it far too easy to start and prolong killing, which ties straight back to *jus ad bellum*, last resort and prospect of success.
- The other side: force protection is a genuine moral good, you have a duty to your own soldiers' lives, and there's nothing noble about getting your people killed just to make the fight 'sporting'.
- My reply: protecting your troops is fine, but the deeper unease isn't really about fairness to the enemy, it's that risk-free killing erodes the seriousness and restraint that should surround taking a human life.
- Concrete probe to mention: Reflective Exercise 10.2 asks whether Australian SAS should remove their body armour to make the fight 'fair', which is obviously absurd, and that absurdity shows the worry isn't really about fairness but about what risk-free killing does to us.
- Closing line: the asymmetry isn't wrong because it's unfair to terrorists, it's worrying because it makes war cheap, easy, and detached from its true gravity.

12. Daniel Byman argues 'Why Drones Work'. What's his case, and is he right?

- My view: Byman makes a strong consequentialist case and on precision he's partly right, but he understates the second-order costs, so 'they work' is too quick.
- Byman's core claim: drones are the most precise weapon we have, they let you kill specific high-value terrorists while minimising both civilian deaths and risk to your own troops, compared to the brutal alternatives like full invasions, occupations, or letting Pakistan's army shell whole regions.
- The force-protection point: you decapitate terrorist leadership, disrupt their planning, and you do it without putting a single soldier's life on the line, which on a pure cost-benefit view looks like an ethical win.
- Why this is appealing: if the realistic alternative to a drone strike is a far bloodier ground operation or no action at all, then the drone genuinely is the least-harm option, and refusing it could cost more innocent lives.
- The other side (Boyle): this ignores strategic blowback, every civilian death radicalises families and communities, breeds new recruits, and corrodes the legitimacy of the government you're propping up, so drones can be counter-productive and create more enemies than they kill.
- My reply / where I land: Byman wins on the narrow tactical question (drones ARE precise relative to a B-52), but Boyle wins on the strategic one, precision at the point of impact doesn't redeem a programme that makes you less safe over a decade.
- Term to drop: this is a classic consequentialist disagreement, both Byman and Boyle judge drones by outcomes, they just disagree about which outcomes (immediate kills vs long-run radicalisation) dominate.
- Closing line: drones 'work' as a weapon, the real question Byman doesn't settle is whether the war they enable works as a strategy.

13. Walk me through the cafe cruise-missile case: you can fire a missile at a group of terror plotters in a London cafe, but it'll kill 12 innocent bystanders, or you can let them escape. How do you reason it?

- Opening move: I wouldn't just blurt yes or no, I'd reason it through with the two key tools, proportionality and the doctrine of double effect, because that's what the case is testing.
- First, double effect: the deaths of the bystanders must be foreseen but NOT intended, not your aim and not your means; you're aiming at the plotters and the 12 deaths are a tragic side-effect, not the point of the strike.
- Second, proportionality: even an unintended side-effect has to be worth it, so I'd weigh the harm prevented (how certain and how catastrophic is the plot?) against the certain harm caused (12 dead innocents).
- The numbers matter but aren't decisive: 12 confirmed innocent deaths to stop a vague, uncertain plot fails proportionality, whereas 12 deaths to stop the certain detonation of a dirty bomb in a city might pass it, so I'd press on how imminent and how serious the threat really is.
- Last resort and necessity: is killing genuinely the only option? A cafe in London means police, arrest, and surveillance are realistically available, which is exactly why this looks more like law enforcement than war, and that pushes me hard toward NOT firing.

- Where I land: in the case as stated, a London cafe with 12 certain civilian deaths and an arrest plausibly on the table, I'd say don't fire, the proportionality and last-resort tests both fail.
- Thinker / contrast to drop: a strict deontologist or absolutist like Nagel would say killing 12 innocents as a side-effect can still be impermissible regardless of the good outcome, while a utilitarian would just run the numbers, naming that clash shows I see the deeper theory fight underneath.
- Closing line: the case is really a test of whether you can show your working, the answer matters less than reasoning aloud through intention, proportionality, and last resort.

14. The August 2021 Kabul drone strike killed ten civilians, including seven children, instead of an ISIS-K bomber. What does that case tell us about the ethics of drone warfare?

- My view: the Kabul strike is the clearest real-world illustration that drone 'precision' is a myth at the level that actually matters, the intelligence, not the missile.
- What happened, briefly: in the chaotic final days of the Afghanistan withdrawal, the US tracked a car they believed carried an ISIS-K suicide bomber, struck it, and it turned out to be an aid worker, Zemari Ahmadi, ferrying water, ten civilians died including seven children.
- The principle in play: this is a failure of distinction and of precautions in attack, even if the weapon hits exactly what it's aimed at, you've murdered the wrong people if your identification is wrong.
- Why it punctures the Byman defence: defenders sell drones as surgically precise, but precision of the warhead is worthless when the 'pattern of life' analysis (jerry cans read as bomb components, normal movements read as hostile) is built on confirmation bias and time pressure.
- It also showcases the asymmetry problem: operators thousands of miles away, under pressure, with grainy footage, made a fatal call about people they never really saw, no human on the ground to sanity-check it.
- The accountability angle: notably, no one was disciplined, which previews Sparrow's responsibility-gap worry from Week 11, even with a human in the loop, real accountability for these killings is thin to non-existent.
- Term to drop: this is a breach of 'precautions in attack', the in-bello duty to take feasible steps to verify your target and avoid civilians, and double effect can't excuse it because the bystanders weren't even a side-effect of hitting a real target, the target itself was innocent.
- Closing line: Kabul shows the hardest problem in drone war isn't the trigger, it's the certainty behind it, and that certainty was catastrophically false.

15. Boyle warns of 'blowback' from drone strikes. Is the blowback argument a knockdown objection, or can a defender of drones answer it?

- My view: blowback is the single strongest objection to the drone programme, but it's an empirical claim, so it's powerful rather than strictly knockdown.
- Boyle's argument: drone strikes, especially ones that kill civilians, radicalise the victims' families and communities, generate propaganda for the enemy, and erode the legitimacy of the local government, so over time you create more terrorists than you eliminate, it's counter-productive.

- Why this bites so hard: it beats the defenders on their own turf, the whole pro-drone case from Byman is consequentialist ('fewer deaths, safer outcomes'), and Boyle says fine, let's count ALL the consequences, including the long-run ones, and now you're losing.
- The defender's reply: blowback is hard to measure, you can't easily prove the strike created more enemies than the bomber it stopped would have killed, and maybe the radicalisation would have happened anyway, so the empirical case is contestable.
- A second defender move: target selection and tighter rules could reduce civilian deaths and therefore reduce blowback, so the objection counts against bad drone practice, not drones as such (this echoes the 'in principle vs in practice' distinction, Dige).
- My reply / where I land: even if blowback can't be proven to the last decimal, the strategic logic is compelling and matches what we saw in Pakistan, the burden should be on the state doing the killing to show it's NOT counter-productive, and it can't.
- Thinker contrast to drop: Byman vs Boyle is the set-piece debate, both consequentialists, disagreeing about whether immediate tactical wins or long-run strategic costs dominate, naming that frames the whole question.
- Closing line: blowback turns the drone advocates' own scoreboard against them, and that's why it's the objection they most need to answer and least can.

16. Defenders say drones are more discriminate and cause less harm than the alternatives, so refusing to use them is the immoral choice. How do you respond?

- My view: this 'lesser evil' argument has real force, but it smuggles in a loaded comparison and ignores costs that don't show up in the immediate body count.
- The argument stated fairly: compared to carpet-bombing, artillery, or a full ground invasion, a drone strike kills far fewer civilians and risks no friendly troops, so if you're going to use force at all, drones are the most humane option and squeamishness about them costs lives.
- Why it's tempting: it's a clean consequentialist case, on the narrow comparison the drone genuinely does minimise harm, and there's a real moral duty not to use more destructive means than necessary (the necessity principle).
- First problem, the false menu: the choice is rarely 'drone strike vs invasion', it's often 'drone strike vs arrest, surveillance, or doing nothing', and once you widen the options, the drone is no longer obviously the least-harm path.
- Second problem, it lowers the threshold of war: precisely because drones are cheap, riskless, and 'clean', they make killing the default first resort rather than a last resort, so the technology expands the total amount of killing even as it makes each strike smaller.
- Third problem, the long run: Boyle's blowback means today's 'fewer deaths' can produce tomorrow's larger conflict, so the lesser-evil maths only works if you stop counting early.
- Term / thinker to drop: this is a pure consequentialist 'minimise harm' argument, and the deontological reply is that some considerations, due process, accountability, the wrongness of executing people outside war, aren't cancelled just by a lower casualty count, which is the same structure as the Arkin vs Sparrow fight in Week 11.

- Closing line: drones may be the lesser evil in a fixed two-option scenario, but their real danger is that they quietly create the very scenarios in which more killing looks reasonable.

Autonomous weapons / killer robots

17. Explain Sparrow's 'responsibility gap' argument against autonomous weapons. Is it convincing?

- Stance: I find the responsibility-gap argument convincing — autonomous weapons should be banned because they kill people while leaving no one who can justly answer for it.
- The core reason: a basic rule of fighting justly (*jus in bello*) is that if a wrongful killing happens, SOMEONE must be able to be held responsible for it — that's part of treating killing with the seriousness it deserves.
- Sparrow's move: with a genuinely autonomous machine that picks its own targets, no one fits the bill — the programmer couldn't predict that specific act, the commander didn't make the decision, and you obviously can't blame or punish the machine itself.
- So there's a 'gap' where responsibility should be — and using such a weapon is therefore illegitimate.
- Strongest objection: maybe we just assign responsibility by rule — say the commander 'owns' whatever the machine does, like a CO is responsible for their troops. Sparrow's reply: that's unfair if the machine genuinely acted in a way the commander couldn't foresee or control — you'd be scapegoating someone for a choice they didn't make.
- Concrete hook: the 2021 Kabul drone strike killed 10 civilians and almost no one was held accountable even WITH a human deciding — now imagine removing the human entirely; the gap gets worse.
- Term/thinker to drop: this is Sparrow's argument from 'Killer Robots' / 'Robotic weapons and the future of war'; the precise label is the 'responsibility gap'.
- Closing line: the deeper point is that responsibility isn't a nice add-on to warfare — it's part of what makes killing in war morally tolerable at all, so a weapon that erases it is wrong even if it's accurate.

18. Arkin says robots could actually be MORE humane than human soldiers. How does Sparrow respond, and who's right?

- Stance: Arkin makes a real point about outcomes, but I side with Sparrow — better statistics don't settle the moral question.
- Arkin's case ('The Case for Ethical Autonomy'): robots feel no fear, anger, revenge or self-preservation, so an 'ethical governor' could follow the laws of war more reliably than scared, exhausted soldiers — and if they kill FEWER civilians, refusing to build them is the immoral choice.
- Why that's powerful: human soldiers commit atrocities — My Lai is the classic case — precisely because of panic and rage that a machine doesn't have.

- Sparrow's reply — this is a consequentialist argument, and his objections are deontological: even if the body-count is lower, it doesn't fill the responsibility gap (no one answers for the killings) and it doesn't restore the respect owed to the person being killed.
- The clean way to frame it: it's the basic consequentialism-versus-deontology fight wearing a robot costume — outcomes versus the rightness of the act and who's accountable.
- My judgement: I lean Sparrow, but honestly Arkin's challenge bites — if machines really did slash civilian deaths, a pure rules-based 'never' starts to look callous, so the responsibility/respect arguments have to carry the weight.
- Term to drop correctly: Arkin's mechanism is the 'ethical governor'; Sparrow and Arkin are the two recognised poles of this debate and actually debated it directly.
- Closing line: the question isn't just 'do robots kill fewer people?' — it's 'is a death decided by a machine the kind of death a human being is owed?'

19. What is Sparrow's 'Robots and Respect' argument? Does being killed by a machine wrong you in a special way?

- Stance: yes — I think being killed by a machine is a distinct wrong, on top of the death itself, because no human agent ever acknowledges the gravity of taking your life.
- The argument (Sparrow, 'Robots and Respect'): persons have a special moral status, and how you're treated can fail to respect that status even if the outcome is 'the same' — a machine that 'decides' your death treats you as a target to be processed, not a human whose life matters.
- Plain version: there's something about a person choosing to kill you — and bearing the weight of it — that recognises you as a fellow human; a machine can't do that, so the killing disrespects you.
- The intuition pump (Sparrow's reflective exercise): would it matter if your brother were killed by an autonomous robot rather than a human-piloted drone? Most people feel it **WOULD** matter — and Sparrow says that feeling tracks a real moral wrong, not just squeamishness.
- Strongest objection: in a chaotic war the dead person never meets their killer anyway — a sniper a mile away or a cruise missile is just as impersonal, so why single out the robot?
- Reply: there's still a human moral agent in the loop with the sniper or missile operator who has decided and is answerable; with full autonomy that agent is gone entirely — the difference is presence of a responsible chooser, not physical distance.
- Term to drop: this rests on 'moral status' / dignity — a deontological idea about respecting persons, distinct from the responsibility-gap argument.
- Closing line: so even a perfectly accurate killer robot could wrong you — not by missing, but by killing you as if no human judgement were owed.

20. Could an autonomous weapon ever satisfy the principles of distinction and proportionality?

- Stance: distinction is maybe achievable one day, but proportionality is the real wall — I doubt a machine can ever do it, and that's a strong reason to keep humans in control.

- Quick definitions to anchor it: distinction means only deliberately targeting combatants and sparing non-combatants; proportionality means a strike's expected civilian harm mustn't outweigh the military advantage.
- Why distinction is hard but conceivable: telling a soldier from a civilian is partly a recognition task — uniforms, weapons, behaviour — and AI is getting good at pattern recognition, so in clean cases a machine might manage it.
- Why proportionality is the killer: it isn't recognition, it's JUDGEMENT — weighing 'how much military advantage' against 'how many civilian lives', and that's a contextual moral call, not something you can reduce to a number in code.
- The objection from the other side: Arkin says a machine without fear or rage might actually apply these rules MORE consistently than a panicking soldier who shoots in scenario-fulfilment.
- Reply: consistency isn't the same as correctness — a machine can reliably apply a bad weighting; and worse, if it gets proportionality wrong, you're back in the responsibility gap with no one to answer for it.
- Real case to name-drop: the Kabul 2021 strike and the café/cruise-missile scenario both turn on a human agonising over whether the bystander deaths were 'worth it' — exactly the judgement we're unsure a machine can make.
- Closing line: distinction is an engineering problem; proportionality is a moral one — and you can't compile a conscience.

21. Does it actually matter, morally, whether a human or a machine pulls the trigger?

- Stance: yes, it matters — and not for a sentimental reason but because the human in the loop is what makes the killing answerable and respectful.
- The reason in plain words: when a person decides to kill, two things are present — someone who can be held responsible afterwards, and someone who recognises the human gravity of what they're doing; a machine deciding strips both away.
- This bundles Sparrow's two arguments: the responsibility gap (no one to blame) and robots-and-respect (the dead person's moral status goes unacknowledged).
- The intuition test: imagine your sibling killed by a 'killer robot' versus by a human operator — most people feel the robot version is worse, and Sparrow argues that reaction is tracking a genuine wrong.
- Strongest objection (the sceptic / consequentialist): the bullet doesn't care what fired it — if the machine is more accurate and kills fewer innocents, fixating on 'who pulled the trigger' is fetishising the human and costing real lives.
- Reply: the point was never accuracy — it's that some decisions are too morally heavy to delegate to a thing that can't be held to account or grasp what it's doing; outcomes don't dissolve that.
- Thinker to drop: Sparrow says yes it matters; a strict consequentialist (the Arkin side) says no, only results matter — so this is the consequentialism-versus-deontology divide in miniature.
- Closing line: if it would horrify you for a machine to decide your brother's death, you already believe the trigger-puller matters.

22. What does 'meaningful human control' mean, and is it the right solution to the killer-robot problem?

- Stance: 'meaningful human control' is the right demand in principle, but it's slippery — and unless it's genuinely meaningful, it just launders the same problems.
- What it means: the policy idea is to keep a human 'in the loop' (approving each target) or 'on the loop' (able to veto), so a real person, not the software, is making the kill decision.
- Why it's attractive: it directly plugs Sparrow's responsibility gap — if a human actually decides, that human is answerable — and it preserves the respect a human agent shows in taking a life.
- Who's pushing it: the Campaign to Stop Killer Robots and the viral 'Slaughterbots' video popularised the demand for a human to stay in control.
- The catch / objection: control can be 'meaningful' on paper but hollow in practice — a tired operator rubber-stamping a machine's recommendation in two seconds isn't really deciding; that's 'automation bias', and it recreates the gap while pretending to close it.
- Reply / where I land: so the test has to be substantive — the human needs real time, real information, and real ability to refuse — not just a button; otherwise it's control in name only.
- Link to drop: this connects to signature strikes from Week 10, where 'targeting by pattern of life' already pushes humans toward deferring to the system's judgement.
- Closing line: 'meaningful human control' is the right flag to plant — the whole fight is over making the word 'meaningful' actually do some work.

23. One worry about autonomous weapons is that they 'lower the threshold of war.' Explain that, and how serious is it?

- Stance: I think this is one of the most serious worries — autonomous weapons make starting wars too easy, and that's a *jus ad bellum* problem, not just a battlefield one.
- The reason: a big political brake on going to war is the cost of sending your own people home in body bags; if your soldiers aren't at risk because robots do the fighting, that brake is released.
- So leaders face less domestic resistance and may reach for force where they'd previously have negotiated — war becomes a cheaper, more tempting tool.
- This links to the asymmetry / 'riskless war' worry from Week 10 drones — Sparrow is sceptical that fighting with zero risk to yourself is the morality-free win it looks like.
- Strongest objection: surely fewer of MY soldiers dying is just good? Force protection is a legitimate aim, and Byman's 'Why Drones Work' argues precision plus safety is a moral plus, not a minus.
- Reply: the harm isn't to your soldiers — it's that lowering the cost of war means MORE wars and more civilian victims overall, plus Boyle's 'blowback' point that easy strikes breed resentment and future conflict.
- Real cases: the proliferation of drones already shows states striking across borders with little political cost — autonomous systems would amplify that further.
- Closing line: a weapon that makes war painless for the side that has it doesn't make war more humane — it just makes it more frequent.

24. Give me the strongest case FOR building autonomous weapons — then tell me whether it survives.

- Stance (steelman first): the strongest pro-case is genuinely strong — if killer robots reliably kill fewer civilians, building them could be a moral duty, not a crime — but I think it ultimately fails on accountability.
 - The pro-argument (Arkin): machines don't feel fear, rage, revenge or self-preservation, so an 'ethical governor' could follow the laws of war more consistently than stressed soldiers who panic — fewer My Lais.
 - Add the consequentialist clincher: if the data showed AWS halved civilian deaths, then refusing to deploy them means choosing MORE dead innocents out of squeamishness — that's a serious charge.
 - Throw in force protection: Byman's 'Why Drones Work' line that keeping your own people out of harm's way while striking precisely is a real ethical benefit.
 - Now the turn — why it doesn't survive for me: Sparrow's two replies stand untouched by the body-count — the responsibility gap means wrongful killings have no one to answer for them, and robots-and-respect means even 'correct' kills disrespect the victim's moral status.
 - Plus the proportionality doubt: 'fewer civilians' assumes the machine can weigh military advantage against civilian harm — but that's judgement, not recognition, and I don't think code can do it.
 - And the threshold worry: even if each strike is cleaner, making war cheaper to wage likely means MORE war overall, so the civilian-saving math may not hold at the strategic level.
 - Closing line: the pro-case wins the spreadsheet but loses the argument — accuracy can't supply accountability or respect, and those are what make killing in war bearable at all.
-

Q2 — Foundations (weeks 1–8)

One question is on the foundations: can war be right, when, and the limits.

Foundations: relativism, realism, pacifism, just war theory

25. This is an ethics unit, but someone might say morality is just relative — different cultures, different rules, who are we to judge? Why does the unit reject ethical relativism, and do you agree?

- My stance: I think the unit is right to set relativism aside — if we took it seriously we couldn't even get the ethics of war off the ground, so we have to assume some moral answers are genuinely better than others.
- The main reason in plain words: if morality were just opinion, then saying 'a massacre is wrong' would only mean 'I personally dislike massacres' — but that's clearly too weak; we want to say it's actually wrong, full stop.
- The other side: the relativist has a fair point that cultures really do disagree, and there's a worry about arrogantly imposing our values on others.

- My reply: disagreement about morality doesn't prove there's no right answer — people disagree about facts too, and that doesn't make facts relative; we can still reason and argue toward better answers.
- Concrete case: war crimes like Bucha in Ukraine or My Lai — we don't want to say those are just 'wrong-for-us'; we want to say they're wrong period, and relativism can't give us that.
- Term/thinker to drop: the unit's 'anti-relativism ground rule' — and the descriptive-versus-normative distinction; relativism confuses how people DO behave with how they OUGHT to.
- Closing line: so I'd say relativism is a non-starter for this unit — you can't condemn an atrocity if everything is just opinion, and the whole point of the ethics of war is that some things really are off-limits.

26. The realist says morality just doesn't apply to war — states do whatever it takes to survive. Is the realist right?

- My stance: the realist is partly right as a description of how states behave, but wrong to conclude morality is irrelevant — I'd reject the strong version.
- The main reason: realists like Morgenthau argue the world is anarchic, there's no world police, so a state that fights morally while its enemies don't will simply lose — survival comes first.
- The other side / why it's tempting: there's something honest about it — leaders DO act on national interest, and pretending otherwise can be naive.
- My reply: even realists smuggle morality back in — they say a leader 'ought' to protect their people, which is itself a moral claim; and we clearly DO judge states, e.g. we condemn aggression and war crimes.
- Concrete case: we don't just shrug at Russia invading Ukraine or at Bucha and say 'that's just power politics' — we make moral judgements, which shows morality isn't actually irrelevant.
- Term/thinker: Morgenthau and Bacevich for realism; and note realism leans on a hard-nosed consequentialism about state survival — Orend is the just-war theorist who rebuts both realism and pacifism.
- Closing line: so I'd say the realist describes a real pressure, but 'might makes right' isn't a moral theory — it's a refusal to do ethics, and that's exactly what this unit pushes back on.

27. People throw the word 'realism' around loosely. Can you explain the difference between descriptive and prescriptive realism, and why it matters?

- My stance: this distinction is crucial — descriptive realism can be true while prescriptive realism is false, so collapsing them is the big mistake.
- Descriptive realism in plain words: it just says how states ACTUALLY behave — they pursue power and self-interest, morality often takes a back seat. That's an empirical claim, and it's largely accurate.
- Prescriptive realism: it goes further and says states OUGHT to ignore morality — that ethics SHOULDN'T constrain them. That's a normative claim, and it's much more controversial.
- Why it matters: even if you grant the descriptive bit, the prescriptive bit doesn't follow — 'states do behave selfishly' doesn't prove 'states should behave selfishly'; that's the classic is-ought gap.

- The reply to prescriptive realism: it's basically self-defeating, because saying a leader OUGHT to put survival first is itself a moral 'ought' — so you can't escape ethics, you've just picked a thin one.
- Term/thinker: Morgenthau's 'twilight of international morality' and Bacevich; tie it to the descriptive-versus-normative point — this unit is normative, so it targets prescriptive realism.
- Closing line: so when someone says 'be realistic,' I'd ask which realism — the true-but-harmless descriptive one, or the prescriptive one that actually tries to cancel ethics and doesn't succeed.

28. What's the strongest version of pacifism you can give — and what's the best objection to it?

- My stance: I think the strongest pacifism is the absolute, principled version — that killing human beings is simply always wrong — but ultimately I find the objection to it decisive.
- The strongest case in plain words: Tolstoy's line that violence only breeds more violence and that patriotism is the root of war — if every nation renounced force, you'd break the cycle; war is mass killing, and killing is wrong.
- Distinguish two flavours to show range: absolute pacifism (killing is wrong in itself, a deontological view) versus contingent pacifism (the costs of war always outweigh the benefits, a consequentialist view).
- The best objection — Anscombe: in 'War and Murder' she calls pacifism 'pernicious' because it blurs the line between killing combatants and murdering the innocent — if all killing is equally forbidden, you lose the crucial distinction that protects civilians.
- Why the objection bites: if killing an armed aggressor is 'just as bad' as killing a baby, then a soldier who refuses to fire has no special reason to spare non-combatants — pacifism accidentally undermines the rule it should care most about.
- Concrete case: a defensive war against an invader, or stopping a genocide — refusing all force can mean allowing a massacre, which looks like a worse moral outcome.
- Thinker note used correctly: Anscombe is NOT a pacifist — she's a just-war Catholic who defends the Doctrine of Double Effect; she attacks pacifism precisely to PROTECT non-combatant immunity.
- Closing line: so the strongest pacifism is morally serious, but Anscombe shows it proves too much — by forbidding all killing equally, it weakens the very protection of the innocent that should be its whole point.

29. So if realism and pacifism are both out, what is just war theory — and why is it the unit's main framework?

- My stance: just war theory is the sensible middle path — war is neither always fine nor always forbidden; it's sometimes justified, but only if strict conditions are met.
- The main idea in plain words: it threads the needle between realism ('anything goes') and pacifism ('nothing goes') by saying war can be morally OK, but it's heavily constrained by rules.

- The structure to name: it splits into *jus ad bellum* — the justice of GOING to war — and *jus in bello* — justice IN how you fight; and these are independent, so you can fight a just cause unjustly, or vice versa.
- The *ad bellum* criteria to rattle off: just cause, legitimate authority, right intention, proportionality, last resort, and reasonable prospect of success.
- The *in bello* core: discrimination — only target combatants, so non-combatant immunity — plus proportionality and not using methods that are evil in themselves (*mala in se*).
- Why it's philosophically rich: it blends both master theories — deontological limits (never deliberately target the innocent) WITH consequentialist checks (is the good worth the cost, i.e. proportionality).
- Concrete case / thinker: Orend defends just war theory against both realism and pacifism; and you can test it on a live case like the US/Israel strikes on Iran — was there just cause, was it last resort, was it proportionate?
- Closing line: so I'd say JWT is the unit's backbone because it takes morality seriously WITHOUT being naive — it lets us actually judge a war as just or unjust instead of either excusing everything or condemning everything.

30. Just war theory tries to sit between realism and pacifism — but is it really a stable middle, or does it just borrow from both and please no one?

- My stance: I think it IS a stable and defensible middle — it's not fence-sitting, it's a principled combination, and that's a strength rather than a weakness.
- The main reason: it answers the one question both extremes get wrong — realism says morality never applies, pacifism says it always forbids; JWT says morality applies AND can sometimes permit, which matches our actual moral judgements.
- The objection from the realist side: they'd say the rules are unrealistic — in a real war, states will break them to survive, so JWT is just nice-sounding theory.
- The objection from the pacifist side: they'd say once you permit ANY war, you open the door to endless justified killing — the criteria are too easy to fudge.
- My reply: the fact that rules get broken doesn't make them wrong — we still use them to JUDGE and condemn, e.g. the Brereton report into alleged Australian war crimes used *in-bello* standards; and the criteria are demanding, not a rubber stamp.
- Term/thinker: Orend explicitly defends JWT by rebutting both realism and pacifism; and the deontology-plus-consequentialism blend is what lets it borrow the best of each without collapsing into either.
- Closing line: so I'd say it's not 'pleasing no one' — it's the only view that lets us say THIS war was just and THAT one wasn't, which is exactly what doing the ethics of war requires.

31. Suppose a friend says 'all this is just armchair philosophy — in a real war none of these moral rules survive contact with the enemy.' How would you defend the whole project of the ethics of war?

- My stance: I'd push back hard — the ethics of war isn't armchair stuff; it's exactly what we reach for when we judge real atrocities, so it's deeply practical.
- The main reason: we obviously DO make these judgements — we call some acts war crimes, we prosecute them, we condemn aggression — and you can only do that if there are real moral standards, not just opinions.
- The descriptive-versus-normative move: my friend is sliding from 'rules get broken' (descriptive) to 'rules don't matter' (normative) — but laws against theft get broken too, and we don't conclude theft is fine.
- The realist objection acknowledged: yes, states under pressure cut corners — that's the grain of truth in Morgenthau's realism — but 'people misbehave' is a reason to HAVE standards, not to abandon them.
- Why the standards bite in practice: international law, the Geneva Conventions, and tribunals all encode just-war ideas like distinction and proportionality — so the 'armchair' rules literally shape real prosecutions.
- Concrete case: My Lai, Bucha, or the Brereton report — in each one we don't say 'war is war'; we say a line was crossed, which only makes sense if the line is real.
- Term/thinker: Orend's rebuttal of realism is the textbook answer here — plus the anti-relativism ground rule, since both relativism and crude realism end up saying 'no real wrongs exist.'
- Closing line: so I'd tell my friend the rules surviving 'contact with the enemy' isn't the test — the test is whether we can still call an atrocity an atrocity, and we plainly can, which is the whole point of doing this.

Starting a war + humanitarian intervention

32. Is going to war in national self-defence morally the same as an individual defending themselves from attack? (Walzer vs Rodin)

- STANCE: I'd say it's NOT as simple as Walzer makes it sound — the analogy between a nation defending itself and a person defending themselves is shakier than it looks, and Rodin is right to press on it.
- THE STANDARD VIEW (Walzer's 'legalist paradigm'): a state is like a person — aggression against its sovereignty or territory is like someone attacking you in the street, so the state has a clear right of self-defence, and that's what justifies going to war.
- RODIN'S CHALLENGE ('the reductivist problem'): if national self-defence just reduces to a bunch of individuals defending their lives, it can't justify everything we think war can — e.g. fighting to the death over a border or territory when surrender would actually save more lives.
- WHY THAT BITES: in personal self-defence you can only use lethal force to stop a lethal threat to a PERSON — but states often fight for things like territory, autonomy, or 'national honour', which aren't individual lives, so the analogy breaks.
- REPLY / why self-defence still matters: you can say states are MORE than the sum of individuals — there's a 'common life' or shared political community worth defending, so defending the state isn't just defending bodies; but admit this needs an extra argument Walzer kind of assumes.

- TERM/THINKER to drop: Walzer's 'domestic analogy' / 'legalist paradigm'; Rodin's 'reductivist challenge' — and note this is the same revisionist move McMahan makes about combatants.
- CLOSING LINE: so I'd land on — self-defence is a real just cause, but Rodin shows we can't just lazily copy the individual case across; the state version needs its own justification.

33. Walk me through the jus ad bellum criteria — what has to be true for it to be just to GO to war?

- STANCE: jus ad bellum is the checklist for whether STARTING a war is justified — and the key point is you need ALL the boxes ticked, not just one, otherwise it's unjust even if your cause is good.
- JUST CAUSE: there has to be a genuine wrong to respond to — classically defence against aggression; the textbook hard case is whether protecting foreigners (humanitarian intervention) counts.
- LEGITIMATE AUTHORITY + RIGHT INTENTION: it has to be declared by a proper authority (a state, arguably the UN), and you have to actually be fighting FOR the just cause — not using it as a cover for grabbing oil or territory.
- PROPORTIONALITY + LAST RESORT: the overall good has to outweigh the harm and destruction (this is Hurka's territory — what goods even count?), and you must have genuinely tried or exhausted peaceful options first.
- REASONABLE PROSPECT OF SUCCESS: there's no point starting a war you'll lose — pointlessly spending lives is itself wrong, so a hopeless war fails the test even with a just cause.
- WHY IT'S A PACKAGE: I'd stress these work TOGETHER — a war can have a just cause but fail on proportionality or last resort, and then it's still unjust; that's what makes it a real test, not a rubber stamp.
- EXAMPLE/THINKER: I'd name-check the 2003 Iraq case for failed last resort / dodgy intention, or Libya 2011 for a war that arguably ticked the boxes going in but went wrong; Hurka as the proportionality guy.
- CLOSING: so going to war is presumptively wrong and the criteria are hurdles you have to clear — the burden is on the one starting it.

34. What is proportionality in jus ad bellum, and why is it so hard to actually apply? (Hurka)

- STANCE: proportionality sounds simple — the good of the war has to outweigh the harm — but Hurka shows it's genuinely messy because you have to decide WHICH goods and harms even count, and that's a judgement call.
- THE BASIC IDEA: you weigh the good you're trying to achieve (stopping aggression, saving lives) against the destruction, deaths and suffering the war will cause — if the harm is bigger, the war fails even if the cause is just.
- WHY IT'S HARD (Hurka's point): which goods go on the scale? Just stopping the aggression, or also things like deterring future aggressors, or 'national honour'? Hurka argues only some goods are 'relevant' and you can't just throw everything in.

- **ANOTHER DIFFICULTY:** you're weighing things that don't share a unit — lives against freedom against future risks — and you're doing it in advance, under huge uncertainty about how the war will actually go.
- **OBJECTION:** a critic says this makes proportionality hopelessly vague — if you can't measure it, it's useless as a criterion.
- **REPLY:** even rough, it still rules out the obvious cases — a massive war over a tiny grievance clearly fails — so it does real work as a constraint even without a precise formula.
- **TERM/THINKER:** Hurka on 'relevant goods'; note proportionality appears in BOTH halves — jus ad bellum (the whole war) and jus in bello (each strike).
- **CLOSING:** so proportionality is less a calculator and more a discipline that forces you to justify the costs honestly.

35. Does the 'Responsibility to Protect' (R2P) override a state's sovereignty? (Annan vs Coady)

- **STANCE:** I lean towards R2P having real force — sovereignty shouldn't be a shield for mass atrocity — but I take Coady's worries seriously enough that I'd want strict limits.
- **ANNAN'S CASE:** 'sovereignty as responsibility' — a state earns its sovereignty by protecting its own people, so if it's slaughtering them (genocide, ethnic cleansing), it forfeits the right to be left alone, and the international community can step in.
- **THE CORE TENSION:** this pits two values against each other — sovereignty / non-interference (the bedrock of international order) versus universal human rights that don't stop at borders.
- **COADY'S CRITIQUE:** humanitarian war is dangerous — it can be a fig leaf for self-interested intervention, it often makes things worse, and 'we're saving them' is exactly what aggressors always say; good intentions don't guarantee good outcomes.
- **KRASNER ANGLE:** sovereignty is already 'organised hypocrisy' — powerful states break it when it suits them and invoke it when it doesn't — so R2P risks being selectively applied to the weak.
- **REPLY / why the stance holds:** the answer to abuse isn't to abandon R2P but to constrain it with the jus ad bellum criteria — legitimate authority (UN), last resort, reasonable prospect of success — so it can't just be unilateral cover.
- **EXAMPLE/THINKER:** Libya 2011 as the test case — started as R2P, drifted into regime change (Kuperman's reassessment); Rwanda as the case where NOT intervening was the scandal.
- **CLOSING:** so R2P can override sovereignty, but only as a tightly fenced exception, not a general licence.

36. Was NATO's 2011 intervention in Libya a success or a cautionary tale for humanitarian war?

- **STANCE:** I'd call Libya the cautionary tale that humanitarian-war sceptics point to — it shows how a just-cause intervention can satisfy the criteria going in and still end badly.
- **THE CASE FOR IT (going in):** Gaddafi was threatening a massacre in Benghazi, there was a UN Security Council resolution (legitimate authority), and the stated aim was protecting civilians — textbook R2P.

- **WHAT WENT WRONG:** the mission slid from 'protect civilians' into regime change and toppling Gaddafi, which went beyond the mandate; afterwards Libya collapsed into chaos, civil war and a failed-state vacuum.
- **KUPERMAN'S REASSESSMENT:** he argues the original 'imminent massacre' threat was overstated, and the intervention actually prolonged and worsened the conflict — so even the good-cause premise is questionable.
- **WHY IT MATTERS FOR R2P:** it's exactly Coady's fear realised — the gap between intention and outcome, and 'mission creep' from protection to regime change, which poisoned appetite for later interventions (e.g. Syria).
- **OBJECTION / the other side:** a defender says doing nothing might have meant a Benghazi massacre, and the post-war collapse was a failure of follow-through, not of the decision to intervene.
- **REPLY:** but 'reasonable prospect of success' and proportionality are jus ad bellum criteria — if you can't secure the peace afterwards, that counts against going in; intentions aren't enough.
- **CLOSING:** so Libya is the worked example of why R2P needs the full jus ad bellum discipline, not just a noble motive.

37. Is sovereignty something we should really respect, or is it just 'organised hypocrisy'? (Krasner)

- **STANCE:** I'd say sovereignty matters as a principle of order, but Krasner's right that in practice it's applied so inconsistently it's partly a convenient fiction.
- **WHAT SOVEREIGNTY IS:** the idea that a state runs its own affairs free from outside force — non-interference is the default rule of the international system and what aggression violates.
- **KRASNER'S 'ORGANISED HYPOCRISY':** powerful states constantly breach others' sovereignty when it suits them and loudly defend it when it doesn't — so the norm is honoured selectively, not consistently.
- **WHY THIS BITES ON HUMANITARIAN WAR:** if sovereignty is selectively enforced, then R2P interventions risk being applied only to weak states the powerful want to act against — never to allies doing similar things.
- **THE OTHER SIDE:** you can defend sovereignty as still valuable even if imperfectly applied — it protects weak states from strong ones and gives smaller nations a seat at the table; hypocrisy doesn't make the rule worthless.
- **REPLY / balance:** so I'd hold both — sovereignty earns respect as protection against aggression (Walzer's reason), but Krasner stops us treating it as sacred when it's shielding atrocities.
- **THINKER/LINK:** tie Krasner to Annan's 'sovereignty as responsibility' — both agree sovereignty is conditional, not absolute, just from different angles.
- **CLOSING:** so respect sovereignty, but with eyes open about who actually gets to break it.

38. Reflective one: is your nationality something you deserve — and does that change how you think about intervening to protect foreigners?

- **STANCE:** honestly, no — where I was born is pure luck, not something I earned, and once you admit that it gets much harder to say foreigners' suffering matters less than ours.
- **THE CORE INTUITION:** nobody chooses their birthplace; a person born in Australia didn't do anything to deserve safety any more than someone born in a war zone deserved the war — it's a 'moral lottery'.
- **WHY IT BEARS ON INTERVENTION:** if nationality is morally arbitrary, then the basic interests of foreigners — not being massacred — count just as much as ours, which pushes towards taking R2P seriously.
- **THE OTHER SIDE:** but there's a real worry about partiality — we DO have special duties to our own (family, fellow citizens), and a government's first job is its own people, so 'everyone counts equally' can't be the whole story.
- **REPLY:** special duties can coexist with basic universal duties — I can favour my own AND still owe strangers a floor of 'don't let them be slaughtered if I cheaply can'; arbitrariness undercuts indifference, not all partiality.
- **LINK TO SOVEREIGNTY:** this is why Annan's 'sovereignty as responsibility' is appealing — borders are partly arbitrary lines, so they shouldn't be an absolute moral wall against helping people on the other side.
- **CAUTION (Coady):** but caring about foreigners doesn't license reckless war in their name — good intentions plus arbitrary luck still have to clear proportionality and prospect of success.
- **CLOSING:** so the arbitrariness of nationality doesn't force open borders or constant intervention — it just stops me pretending strangers' lives weigh less, which is enough to take humanitarian duties seriously.

How you fight: distinction, terrorism, evil weapons

39. What makes someone a legitimate target in war? Is it the uniform they wear or the threat they pose?

- My stance: I lean towards McMahan's view that it's about the threat you pose and your responsibility for it, not just the uniform you happen to be wearing.
- The plain reason: distinction (also called discrimination) says you can only deliberately kill combatants, not civilians, because combatants are the ones doing the fighting and posing the threat.
- The traditional view, from Walzer, is simpler: anyone in uniform is fair game, because by becoming a soldier you've made yourself 'dangerous' and waived your right not to be attacked, regardless of which side is just.
- The other side I'd push back on: that uniform-based test treats a conscript who was forced to fight the same as a willing aggressor, which feels wrong because they're not equally responsible.
- McMahan's reply (revisionism): what really makes you liable to be killed is moral responsibility for an unjust threat, so a soldier fighting an unjust war is liable and a soldier fighting a just war is not, even though both wear uniforms.
- Concrete example: a cook or a clearly surrendering soldier still wears the uniform but poses no threat, so the uniform alone can't be the whole story.

- Term to drop: this is the difference between the 'moral equality of combatants' (Walzer) and 'liability' (McMahan).
- Closing line: I'd say the uniform is a useful rough guide in the fog of war, but morally what justifies killing someone is the threat and the responsibility behind it.

40. Why are civilians immune from attack? Can you give me more than one reason?

- My stance: civilian immunity is one of the firmest rules in the ethics of war, and there are at least three separate grounds for it, which is why it's so strong.
- First ground - they pose no threat: civilians aren't fighting you, so on a self-defence logic there's no threat to neutralise, and killing them isn't defending yourself against anything.
- Second ground - they're not responsible: even in an aggressor nation, ordinary civilians didn't start the war and can't stop it, so they're not liable in McMahan's sense of being responsible for an unjust threat.
- Third ground - their moral status / they're 'innocent': in the original sense 'innocent' means 'not harming' (from Latin nocere, to harm), and respect for persons means you can't just use them as a means to your military ends.
- The other side: a consequentialist might say if killing some civilians ends the war faster and saves more lives overall, immunity should bend - this is the utilitarian pressure Nagel writes about.
- The reply: that's exactly what absolutism resists - Nagel's point in 'War and Massacre' is that some things you simply may not do to a person no matter the numbers, because it fails to treat them as a person.
- Example: bombing a residential suburb to break enemy morale (like area bombing in WWII) deliberately targets civilians, and we rightly call that a war crime.
- Closing line: because immunity rests on three independent legs - no threat, no responsibility, and respect for persons - knocking out one doesn't bring the whole thing down.

41. Explain the Doctrine of Double Effect. Does it really get attackers off the hook for civilian deaths?

- My stance: double effect is a genuinely useful moral tool, but it's not a blank cheque - it only works if you take its conditions seriously.
- The core idea in plain words: there's a moral difference between aiming at civilian deaths and merely foreseeing them as a side-effect of hitting a legitimate target, even when the same people die.
- The conditions: you can't intend the civilian deaths as your aim, you can't use them as your means, and the harm has to be proportionate to the military good you're getting.
- The strongest objection: critics say it's too easy to game - you just say in your head 'I didn't intend it, it was foreseen' and suddenly any collateral damage is fine, so it launders bad acts.
- The reply: that's why the proportionality and 'not a means' conditions matter, and why thinkers like Walzer add 'due care' - you must actually accept risk to your own forces to reduce civilian harm, not just shrug.

- Example: a cruise missile on a genuine military target that kills bystanders next door can pass double effect; deliberately shelling a marketplace to terrorise the population cannot, because there the deaths ARE the point.
- Terms and thinkers to drop: intention versus foresight, and the doctrine is defended by Anscombe and discussed by Nagel.
- Closing line: so it doesn't get you off the hook automatically - it draws a line, and plenty of real strikes fall on the wrong side of it.

42. How would you define terrorism? What's the key thing that makes an act terrorism rather than just ordinary warfare?

- My stance: I'd follow Primoratz - terrorism is the deliberate use of violence against innocent people in order to coerce or intimidate some other group into doing what you want.
- The key feature: it's the deliberate targeting of the innocent - the victims aren't the real target, they're a means to pressure a government or population through fear.
- Why that definition is good: it's morally neutral about who does it, so it doesn't just mean 'violence by people I dislike' - it focuses on the structure of the act, not the cause.
- The other side: some say one person's terrorist is another's freedom fighter, so the word is just propaganda - any definition is biased.
- The reply: Primoratz answers that even a just cause doesn't license deliberately killing innocents, so freedom fighters can be terrorists too - the cause and the method are separate questions.
- Notice it tracks just war theory: terrorism is basically a violation of distinction or non-combatant immunity, the same rule that constrains armies.
- Example: a bombing in a cafe to pressure a government is terrorism whether the bombers' grievance is just or not, because the dead bystanders are being used to send a message.
- Closing line: so the heart of it isn't the politics or the size of the bang - it's that the innocent are deliberately attacked to frighten someone else.

43. Is terrorism morally worse than conventional war? Or should we judge them by the same standard?

- My stance: I'd side with Held - we should judge terrorism by exactly the same moral standard we use for war, not a special harsher one.
- The plain reason: if deliberately killing innocents is what's wrong with terrorism, then states do that too - through area bombing, sieges, or 'collateral' deaths they could have avoided - so consistency demands one standard.
- The common view I'm resisting: that terrorism is uniquely evil, a category apart, while what states do in war is somehow cleaner because it's official and uniformed.
- Held's challenge: that's a double standard - if we condemn the rebel who kills 20 civilians, we can't wave through the state that kills 200 civilians and calls it unfortunate side-effects.
- The reply to Held (for balance): defenders of the distinction say there really is a difference between intending civilian deaths and foreseeing them - double effect again - so terrorism's deliberateness is worse.

- My counter: but states often foresee civilian deaths with near-certainty and proceed anyway, which starts to look a lot like intending them, so the gap narrows.
- Example: compare a suicide bombing in a market with a deliberate WWII firebombing of a city - both deliberately kill civilians to break morale, so why call only one terrorism?
- Closing line: I think honesty requires one yardstick - if it's wrong for the weak to deliberately kill the innocent, it's wrong for the strong too.

44. Can a state commit terrorism? What would state terrorism look like?

- My stance: yes, absolutely - there's no reason terrorism has to be the work of small groups; states can and do use terror as a deliberate tool.
- The plain reason: if terrorism is defined by its structure - deliberately attacking innocents to coerce or intimidate, per Primoratz - then a government doing exactly that is committing terrorism, full stop.
- Why people resist this: we tend to picture terrorists as non-state actors with bombs, and states control the language, so they call their own version 'security operations' or 'collateral damage'.
- The argument for it: deliberately bombing civilian areas to break a population's will to resist is terrorism whether it's done by a militia or an air force - the uniform doesn't change the act.
- Held's point applies here: insisting on one standard means we can't exempt states just because they're states - state terror should be named as terror.
- Example: WWII area bombing of cities to crush morale, or a regime shelling its own civilians to intimidate the opposition, fit the definition cleanly.
- The other side: some reserve 'terrorism' for non-state actors and call state versions 'war crimes' instead - but that's arguably just a labelling choice that lets states off lightly.
- Closing line: I'd say if the act is the deliberate terrorising of innocents, it's terrorism regardless of who's wearing the suit - the state badge doesn't sanitise it.

45. What makes a weapon evil in itself - mala in se? Is it the weapon, or how it's used?

- My stance: a weapon is mala in se - evil in itself - when it can't be used without violating the basic rules of fighting, no matter how careful you try to be.
- The plain reason: the in bello rules require you to discriminate between combatants and civilians and to be proportionate, so a weapon that can't do that - that inevitably hits the innocent - is wrong by its very nature.
- The deeper point: 'mala in se' is a deontological verdict, not a body count - it says some means are off-limits in themselves, echoing Nagel's absolutism that there are things you simply must not do to people.
- The other side: a consequentialist will say no weapon is evil in itself, it all depends on outcomes - even a terrible weapon could be justified if it saved more lives on balance.
- The reply: Nagel's response in 'War and Massacre' is that treating a person as a mere target, or using a weapon that can't even aim at the guilty, fails to respect them as a person - and no good outcome buys that back.

- Example: a nuclear weapon is the classic case - its blast and radiation can't distinguish soldier from child and the harm is wildly disproportionate, so many argue it's mala in se; chemical and biological weapons are similar.
- Term to drop: this connects to 'no mala in se' as a rule of jus in bello, alongside distinction and proportionality.
- Closing line: so it's not just how you use it - some weapons are built in a way that makes lawful, respectful use impossible, and that's what makes them evil in themselves.

46. If a weapon could be used precisely against only combatants, would that settle whether it's acceptable? Think about nuclear weapons and the deterrence puzzle.

- My stance: precision helps a lot, but it doesn't fully settle it - with nuclear weapons there's a second, weirder problem about merely threatening to use them.
- The first point: if a weapon genuinely could hit only combatants and be proportionate, it would pass distinction and proportionality, so it wouldn't be mala in se - the problem with nukes is they can't do that.
- Why nukes specifically fail: the blast radius and lingering radiation can't distinguish a soldier from a baby, and the scale is so vast that proportionality is basically impossible.
- The deeper puzzle - the deterrence paradox: deterrence works by sincerely intending to do something - massacre civilians in retaliation - that it would be deeply wrong to actually do, so is it ethical to threaten what's unethical to carry out?
- The other side: defenders say if the threat keeps the peace and you never fire, the good outcome justifies the conditional intention - and Hiroshima arguably ended a war and saved lives.
- The reply: Ward Wilson, in 'The Gordian Knot', doubts deterrence even works and argues it may have been the Soviet invasion, not Hiroshima, that ended WWII - so the supposed benefit is shaky.
- Extra angle to mention: Schlosser on accident risk - near-misses mean even a 'never used' arsenal threatens catastrophe - and Glover's moral-history lens on how we drifted into accepting this.
- Closing line: so precision is necessary but not sufficient - even setting aside the blast, holding a city hostage with a threat you couldn't morally honour is its own ethical knot.

Q3 — Synthesis

One question makes you link a foundation idea (weeks 1–8) to a modern case (weeks 9–11).

47. Take the principle of distinction from jus in bello and apply it to autonomous weapons. Can a killer robot actually 'distinguish' combatants from civilians the way the laws of war require?

- STANCE: I'd say distinction is exactly where autonomous weapons run into deep trouble — recognising a target is not the same as judging one, and the law demands judgement.

- **MAIN REASON:** distinction means only deliberately targeting combatants; a machine can do object-recognition (this shape is a person holding a rifle) but distinction in a messy battlefield needs context — is that a soldier surrendering, a farmer with a tool, a child?
- **OTHER SIDE:** Arkin says a machine could actually do distinction BETTER — no fear, no rage, no self-preservation, no 'scenario fulfilment' that makes a scared soldier shoot first, so it might wrongly kill fewer civilians than a human.
- **REPLY:** even granting Arkin on recognition, the in-bello demand isn't just 'hit fewer civilians' — it requires moral judgement about ambiguous cases, and code weighing a surrender or an ambiguous gesture is exactly what we don't trust it to do.
- **EXAMPLE:** think My Lai for the human failure Arkin points to, versus the Kabul-type ambiguity where a man loading water jugs was read as loading bombs — would a robot have read that pattern any better?
- **TERM/THINKER:** this is the principle of distinction (jus in bello) colliding with Sparrow's scepticism — and it links to proportionality too, which needs even more judgement than distinction.
- **CLOSING:** so my view is recognition isn't distinction; until a machine can exercise genuine moral judgement, it can't satisfy the principle that's supposed to protect civilians.

48. Apply the Doctrine of Double Effect to a drone strike that kills civilian bystanders. Does double effect let the operator off the hook for collateral deaths?

- **STANCE:** Double effect can in principle justify some civilian deaths in a drone strike, but it's not a free pass — it only works if you genuinely don't intend the deaths and the strike is proportionate.
- **MAIN REASON:** DDE says you may cause civilian deaths you FORESEE but don't INTEND, as a side-effect of a legitimate strike — so a drone killing a real combatant, with bystanders dying as an unwanted side-effect, can clear the bar.
- **OTHER SIDE:** critics say in practice this gets abused — operators 'foresee' huge collateral counts and just shrug it off as side-effect, which makes DDE a rubber stamp rather than a real constraint.
- **REPLY:** the proportionality clause is the teeth — even if the deaths are unintended, if 12 bystanders die to get one plotter, the military gain doesn't outweigh the harm and the strike fails DDE anyway.
- **EXAMPLE:** walk the café case — fire a cruise missile at terror plotters in a London café and kill 12 innocents, or let them escape? You reason it out: real combatant target (intention ok), but are 12 deaths proportionate, and was it last resort?
- **TERM/THINKER:** DDE rests on the intention-vs-foresight distinction (Anscombe, Nagel) — a deontological tool that lets just-war theory allow SOME collateral without abandoning civilian immunity.
- **CLOSING:** so DDE is a real test, not an excuse — it permits the café strike only if the deaths are truly unintended AND proportionate, and often that second condition is what sinks it.

49. Realism says morality doesn't really apply to war. Apply that to nuclear deterrence — is the bomb just realism vindicated, or does morality still have a grip?

- **STANCE:** I think nuclear deterrence looks like the realist's home turf, but it actually shows morality still has its claws in — because we keep asking whether threatening mass death is right, not just whether it works.
- **MAIN REASON:** realism (Morgenthau, Bacevich) says states pursue survival in an anarchic world, and the bomb is the ultimate survival tool — MAD keeps the peace through self-interest, no morality needed.
- **OTHER SIDE:** but even hardened strategists won't say 'anything goes' — the whole deterrence debate is shot through with moral words like 'unthinkable', and we treat deliberately incinerating cities as a moral horror, not just a policy cost.
- **REPLY:** realism is better read as **DESCRIPTIVE** (how states behave) than **PRESCRIPTIVE** (what they ought to do) — Orend rebuts the leap from 'states act on interest' to 'morality is irrelevant', and the deterrence paradox is a purely moral puzzle realism can't dissolve.
- **EXAMPLE:** the US/Israel-Iran situation in 2026 shows states reaching for deterrence logic, but the global condemnation language proves we still judge it morally, not just strategically.
- **TERM/THINKER:** realism leans on a hard-nosed consequentialism about state survival — Morgenthau's 'twilight of international morality' — but Ward Wilson even doubts deterrence **WORKS**, which knocks out the consequentialist foundation realists assume.
- **CLOSING:** so my view is deterrence isn't proof morality is irrelevant — it's the case where the moral question (is threatening mass murder ever right?) refuses to go away.

50. The nuclear deterrence paradox is supposed to be a puzzle about intention versus action. Use the intention/foresight distinction from your foundations to unpack it — can it be ethical to THREATEN what it'd be wrong to DO?

- **STANCE:** My view is the deterrence paradox is genuinely hard, but I lean toward saying you can't cleanly separate the threat from the act — sincerely intending to do something monstrous is itself morally compromised.
- **MAIN REASON:** deterrence only works if the threat is **SINCERE** — you must really intend to retaliate by incinerating cities — so this isn't a bluff; it's forming a conditional intention to do mass murder.
- **OTHER SIDE:** the defender says intending isn't doing — if the threat succeeds, no one dies, so a conditional intention you hope never to act on, that saves millions by keeping the peace, is surely permissible.
- **REPLY:** but deontology says intentions matter in themselves — if actually launching is murder of innocents (wrong in itself), then sincerely intending it conditionally seems to corrupt you too; you've made yourself ready to commit an atrocity.
- **EXAMPLE:** this is Reflective Exercise 9.2 directly — 'can it be ethical to threaten to do what it would be unethical to do?' — and the conditional-intention problem is Kavka/Gauthier territory.
- **TERM/THINKER:** this is the intention-vs-act puzzle (intending vs doing — the conditional-intention problem, Kavka/Gauthier) at full stretch — a consequentialist (outcomes only) says the threat is fine because it prevents deaths; a deontologist worries the intention itself is the wrong.

- CLOSING: so I'd argue the paradox exposes the consequentialism/deontology split perfectly — if only results count, threaten away; if intentions can be wrong in themselves, deterrence stains you even when it works.

51. McMahan grounds who may be killed in liability — moral responsibility for an unjust threat, not the uniform. Apply that to signature drone strikes, where targets are picked by 'pattern of life' rather than identity.

- STANCE: I think McMahan's liability standard is actually the sharpest weapon AGAINST signature strikes — if liability tracks responsibility for an unjust threat, you can't be liable just for behaving like a 'pattern'.
- MAIN REASON: McMahan says what makes you a legitimate target isn't a uniform — it's your moral responsibility for posing an unjust threat; so killing requires knowing the person IS responsible.
- OTHER SIDE: signature strikes kill people whose identity you don't even know, on statistical behaviour — military-age males in a suspicious gathering — which is exactly NOT confirmation of individual responsibility.
- REPLY: a defender might say patterns are good evidence of liability in a war zone, but McMahan's whole point is liability is individual and moral, not statistical — you can't infer one person's guilt from a crowd's profile.
- EXAMPLE: the Pakistan drone war Coll documents in 'The Unblinking Stare', and signature strikes that killed unidentified people on pattern — versus the Kabul strike where pattern-reading killed 10 civilians.
- TERM/THINKER: this is liability (McMahan) versus the older 'uniform/combatant' test — and it overlaps the war-vs-execution question, since outside declared war you're supposed to ARREST, not kill, unconfirmed suspects.
- CLOSING: so my view is signature strikes flunk McMahan's test — they kill on the basis of looking dangerous, when liability demands actually BEING responsible for an unjust threat.

52. Walzer says all combatants are moral equals — they may kill each other regardless of whose cause is just. Does that idea even make sense when one 'combatant' is an autonomous killing machine?

- STANCE: I'd argue the moral equality of combatants basically breaks down with killer robots — a machine isn't a combatant who shares the soldier's risks and rights, so the symmetry Walzer relies on collapses.
- MAIN REASON: Walzer's idea is that soldiers on both sides face each other as equals, each with the right to kill and the liability to be killed — a kind of mutual risk and mutual standing.
- OTHER SIDE: you might say a robot is just a tool, like a rifle, so the human commander remains the 'combatant' and moral equality still holds between the two human sides.
- REPLY: but if the machine genuinely selects and engages targets without meaningful human control, the human isn't really the one fighting — and a robot has no moral status to be a 'moral equal', which is exactly Sparrow's respect worry.

- **EXAMPLE:** imagine your brother killed by an autonomous machine versus a human-piloted drone — Sparrow's intuition (Reflective Ex 11.5) is most people feel the robot case is worse, because no human equal faced him.
- **TERM/THINKER:** Walzer's moral equality of combatants versus the revisionists (McMahan, Rodin) who already deny it — and AWS push even further by removing the human from one side of the duel entirely.
- **CLOSING:** so my view is killer robots don't just challenge WHO counts as a moral equal — they remove the very thing (a risk-sharing human agent) that made the idea of combatant equality meaningful.

53. Primoratz defines terrorism as deliberate violence against innocents to coerce a wider group. Use that definition to test a signature drone strike — is the US doing something structurally like terrorism?

- **STANCE:** I'd say a signature strike isn't strictly terrorism by Primoratz's definition, but it sits uncomfortably close, and refusing to see that is too convenient for states.
- **MAIN REASON:** Primoratz says terrorism = DELIBERATELY targeting innocents to terrorise/coerce a wider audience — the key word is 'deliberately', the innocents are the means.
- **OTHER SIDE:** a signature strike supposedly targets combatants/threats, so civilian deaths are (claimed as) unintended side-effects — that's double effect, not deliberate targeting of innocents, so it's not terrorism.
- **REPLY:** but if you strike on 'patterns' knowing you'll kill unidentified people, and the wider effect is to terrorise a whole region into fear of the sky, the line between foreseen and intended starts to blur — and Held argues we should judge state violence by the SAME standard as terrorism, not a softer one.
- **EXAMPLE:** the Pakistan drone campaign — communities living in constant dread of strikes — looks coercive on a population even if each strike claims a specific target; Boyle's blowback point is that it radicalises exactly this way.
- **TERM/THINKER:** Primoratz's responsibility-based account of immunity and his definition, versus Held's 'same standard' move and Sparrow's reply in 'Hands up who wants to die?' on responsibility and immunity.
- **CLOSING:** so my view is the definition saves the drone strike from the 'terrorism' label only via the intention clause — and if we're honest about its coercive effect on civilians, we should at least hold it to the same standard Held demands of terrorists.

54. Sparrow's 'Robots and Respect' says being killed by a machine fails to respect your moral status. Doesn't a riskless drone — a human operator a continent away — raise the same worry? Apply respect to drone asymmetry.

- **STANCE:** I think the respect worry partly applies to riskless drones, but it bites harder on autonomous machines — a remote human still decides your death, even if he risks nothing.
- **MAIN REASON:** Sparrow's point is that taking a human life should involve a moral agent who acknowledges the gravity of it — being killed by a thing that 'decides' your death fails to honour your moral status.

- OTHER SIDE: with a drone there IS a human agent deciding, so respect is arguably preserved — the asymmetry is just about risk, and force-protection (keeping your own people safe) is plainly a good thing.
- REPLY: but the asymmetry adds its own worry — when killing is utterly riskless, the killer never stakes anything, and Sparrow is sceptical that risk-free war respects the seriousness of taking life or keeps soldiers honest about it.
- EXAMPLE: Reflective Ex 10.2 — should the SAS remove body armour to make the fight 'fair'? It probes whether riskless killing is troubling; and 11.5 — brother killed by a robot vs a piloted drone — isolates the pure respect intuition.
- TERM/THINKER: this is moral status/respect feeding Sparrow's two arguments — the respect argument and his scepticism about asymmetric, riskless war.
- CLOSING: so my view is the drone keeps a human in the loop, which preserves SOME respect — but it already strains the idea that killing should cost the killer something, and the robot severs it entirely.

55. Drone and autonomous warfare make war riskless for the attacker. Connect that asymmetry back to jus ad bellum — does making war cheap make it too easy to start?

- STANCE: My view is that riskless warfare is a jus ad bellum problem hiding inside a jus in bello technology — when war costs you no soldiers, the bar to GOING to war quietly drops.
- MAIN REASON: jus ad bellum demands last resort and right intention — but those restraints partly held because war was costly in your own blood; remove that cost and leaders reach for force sooner.
- OTHER SIDE: a defender says cheaper, safer war is GOOD — fewer of your own people die, strikes are precise (Byman), so lowering the cost of just action is no bad thing.
- REPLY: but lowering the THRESHOLD of war means more wars, including unjust ones — and 'last resort' becomes hollow when force is the path of least resistance; the in-bello gain (precision) is bought with an ad-bellum loss (restraint).
- EXAMPLE: the spread of drone strikes outside declared wars, and the 'killer robots' / Slaughterbots worry that armies with no human risk start conflicts a casualty-averse public would once have blocked.
- TERM/THINKER: this links Week 10/11 asymmetry to jus ad bellum criteria (last resort, prospect of success) — and to the 'lowering the threshold of war' critique of AWS; Sparrow is sceptical of riskless war for exactly this reason.
- CLOSING: so my view is the deepest danger of riskless weapons isn't how they're used in a war, but that they make starting wars too easy — an ad bellum failure dressed as an in bello upgrade.

56. Nagel's 'War and Massacre' frames the deep clash as absolutism versus utilitarianism. Show how that exact clash plays out BOTH in nuclear deterrence AND in the Arkin–Sparrow robot debate. Is it the same fight twice?

- STANCE: I'd argue it really is the same fight twice — Nagel's absolutism-vs-utilitarianism (which is just deontology vs consequentialism) is the engine under both the nuclear and the robot debates.

- **MAIN REASON:** Nagel says the utilitarian asks only what produces the best outcome, while the absolutist insists some acts — like deliberately targeting the innocent — are forbidden whatever the payoff.
- **NUCLEAR SIDE:** deterrence is utilitarian (threaten cities to save more lives on balance) versus absolutist (incinerating innocents is murder, so even intending it is wrong) — same structure as the deterrence paradox.
- **ROBOT SIDE:** Arkin is the consequentialist (if robots kill FEWER civilians, build them, refusing is the immoral choice) versus Sparrow the deontologist (good body-counts don't fill the responsibility gap or restore dignity).
- **WHY IT'S ONE FIGHT:** in both, the consequentialist says 'count the bodies and pick the lower number'; the deontologist says 'some things — murdering innocents, killing with no one responsible — are wrong in themselves regardless of the count'.
- **TERM/THINKER:** Nagel 'War and Massacre' (W8) supplies the frame; it's the master key — and the concept hierarchy literally calls the robot debate 'a consequentialism-vs-deontology fight wearing a robot costume'.
- **CLOSING:** so my view is once you see the master key, nuclear and robots stop looking like separate topics — they're the same two moral theories colliding over different weapons.

57. Pacifism says war is never justified. Could a pacifist actually WELCOME autonomous weapons if they reduce killing — or does pacifism still condemn them? Apply the pacifist lens to killer robots.

- **STANCE:** My view is it depends which pacifism — an absolute pacifist condemns AWS like all war, but a contingent 'costs always exceed benefits' pacifist faces a real dilemma if robots genuinely kill fewer.
- **MAIN REASON:** absolute pacifism (a deontological version — killing is always wrong) doesn't care that robots are 'cleaner'; a more humane way to do an impermissible thing is still impermissible.
- **OTHER SIDE:** a contingent pacifist, who opposes war because the costs always outweigh the gains, might in theory welcome a weapon that slashes civilian deaths — that's Arkin's challenge aimed at pacifists too.
- **REPLY:** but even the contingent pacifist can reply that AWS LOWER the threshold of war, so you get more wars overall — net suffering rises even if each war is 'cleaner', so the contingent case against war survives.
- **EXAMPLE:** Tolstoy's absolute renunciation of violence on one side; on the other, the Slaughterbots worry that cheap autonomous killing proliferates conflict rather than reducing it.
- **TERM/THINKER:** Tolstoy (absolute pacifism) versus contingent pacifism; and recall Anscombe's 'War and Murder' attack — she calls pacifism 'pernicious' for blurring killing combatants and murdering the innocent, a charge AWS make newly relevant.
- **CLOSING:** so my view is robots don't rescue war for the pacifist — the absolutist condemns them outright, and the contingent pacifist can point to a MORE warlike world, not a less violent one.

58. For the principle of distinction specifically — which is the bigger threat, a nuclear weapon or an autonomous weapon? Compare the two hard cases against the same foundational

principle.

- **STANCE:** My view is the nuke is the worse offender against distinction in its very nature, while the robot's threat to distinction is about WHO (or what) is doing the judging — different failures, but the nuke is categorically indiscriminate.
- **MAIN REASON:** distinction forbids deliberately or recklessly hitting non-combatants — a strategic nuke physically CAN'T be contained to combatants, so it fails distinction by design, arguably making it mala in se.
- **ROBOT SIDE:** an AWS could in principle be precise — Arkin even says more precise than a panicking soldier — so its problem with distinction isn't blast radius, it's whether a machine can JUDGE ambiguous targets at all.
- **WEIGHING IT:** a precise robot that occasionally misjudges a surrender is bad; a city-buster that vaporises a hundred thousand civilians can't even pretend to discriminate — so on distinction alone, the nuke is the clearer violation.
- **REPLY/NUANCE:** but the robot adds a problem the nuke doesn't — even a 'good' kill leaves a responsibility gap and a respect failure, so it's worse on accountability even if better on blast.
- **TERM/THINKER:** distinction and proportionality plus mala in se — the classic deontological "evil in itself" verdict (Nagel); Sparrow's own contribution is the responsibility gap + respect, a different kind of wrong; Hiroshima is the standing case for indiscriminate effect.
- **CLOSING:** so my view is for distinction itself the nuke wins the 'worse' title because it's inherently indiscriminate — but the robot is uniquely bad in a different currency: no one to answer for the deaths.

59. Mala in se means 'evil in itself' — a deontological verdict. Nuclear weapons are often called mala in se. Could autonomous weapons be mala in se too, or is that the wrong category for them?

- **STANCE:** I'd argue nukes are a much better fit for mala in se than robots — the nuke is wrong in its very effects, whereas the robot's wrong is about agency and accountability, which is a different kind of objection.
- **MAIN REASON:** mala in se (a deontological verdict, needs deontology) means a weapon is wrong INTRINSICALLY, not just for its consequences — torture, poison, rape, and arguably the nuke because it can't help being indiscriminate.
- **OTHER SIDE:** you could call AWS mala in se because killing with no responsible agent, and killing in a way that disrespects moral status, are wrong-in-themselves features — that's Sparrow's flavour of objection.
- **REPLY:** but those are objections to HOW the decision is made (no human answerable, no respect), not to the weapon's intrinsic destructive nature — so 'illegitimate' or 'disrespectful' fits AWS better than the classic mala in se of poison gas.
- **EXAMPLE:** a nuke can never be aimed at only combatants — that's intrinsic; a robot rifle could in principle hit only a soldier, so its wrongness isn't baked into the physics the way the nuke's is.

- TERM/THINKER: mala in se (deontology) — Nagel's absolutism backs calling something evil-in-itself; Sparrow's responsibility gap and 'Robots and Respect' supply a NEW category of intrinsic wrong, not the old effects-based one.
- CLOSING: so my view is nukes are mala in se in the classic effects sense, while robots may be intrinsically wrong in a fresh way — wrong because of who's NOT responsible, not because of what they do to bodies.

60. The jus ad bellum / jus in bello independence rule says a war can be just to start but fought wrongly, or unjust to start but fought cleanly. Use that distinction to analyse the deployment of killer robots — is the AWS debate an in-bello or an ad-bellum question?

- STANCE: My view is the killer-robot debate is sneaky because it straddles BOTH halves — Sparrow's core argument is in-bello, but the deepest danger (lowering the threshold of war) is ad-bellum.
- MAIN REASON: the independence rule says how you START a war and how you FIGHT it are separate moral questions — and AWS raise distinct problems on each side, so you have to keep them apart.
- IN-BELLO SIDE: Sparrow's responsibility gap and the distinction/proportionality worries are all about CONDUCT — can the weapon fight justly, and can anyone be held responsible when it doesn't?
- AD-BELLUM SIDE: but if robots make war riskless and cheap, they erode last resort and right intention — leaders go to war too readily, which is an ad-bellum corruption, not a conduct one.
- WHY IT MATTERS: this means a country could deploy AWS in a war that's perfectly just to START, yet fight it illegitimately (responsibility gap); or fight 'cleanly' with robots in a war it should never have begun — the two failures don't cancel.
- TERM/THINKER: jus ad bellum vs jus in bello independence; Sparrow's responsibility gap and 'Robots and Respect' on the in-bello side; the 'lowering the threshold of war' critique on the ad-bellum side.
- CLOSING: so my view is you can't settle the robot question with one verdict — it fails in-bello on accountability AND threatens ad-bellum on restraint, and the independence rule is what lets you see both.